



PyFR: An Open Source Python Framework for High-Order CFD on Many-Core Platforms

F.Witherden, A. Farrington, P. E.Vincent

Department of Aeronautics
Imperial College London

23rd May 2013

Introduction

- Various types of high-order accurate numerical methods for unstructured grids
- Continuous Galerkin (CG), Discontinuous Galerkin (DG), Spectral Difference (SD), Spectral Volume (SV), Flux Reconstruction (FR)

Introduction

- Various types of high-order accurate numerical methods for unstructured grids
- Continuous Galerkin (CG), Discontinuous Galerkin (DG), Spectral Difference (SD), Spectral Volume (SV), Flux Reconstruction (FR)

Introduction

- FR approach was first proposed by Huynh at NASA Glenn in 2007 [I]
- Based on differential form of governing system
- Unifying, can cast various known + new schemes within a single framework

[I] H.T. Huynh. A flux reconstruction approach to high-order schemes including discontinuous Galerkin methods. AIAA Paper 2007-4079. 2007

Theory

- Consider 1D scalar conservation law

$$\frac{\partial u}{\partial t} + \frac{\partial f}{\partial x} = 0$$

- Divide 1D domain into elements

$$\Omega = \bigcup_{n=0}^{N-1} \Omega_n \quad \bigcap_{n=0}^{N-1} \Omega_n = \emptyset \quad \Omega_n = \{x \mid x_n < x < x_{n+1}\}$$



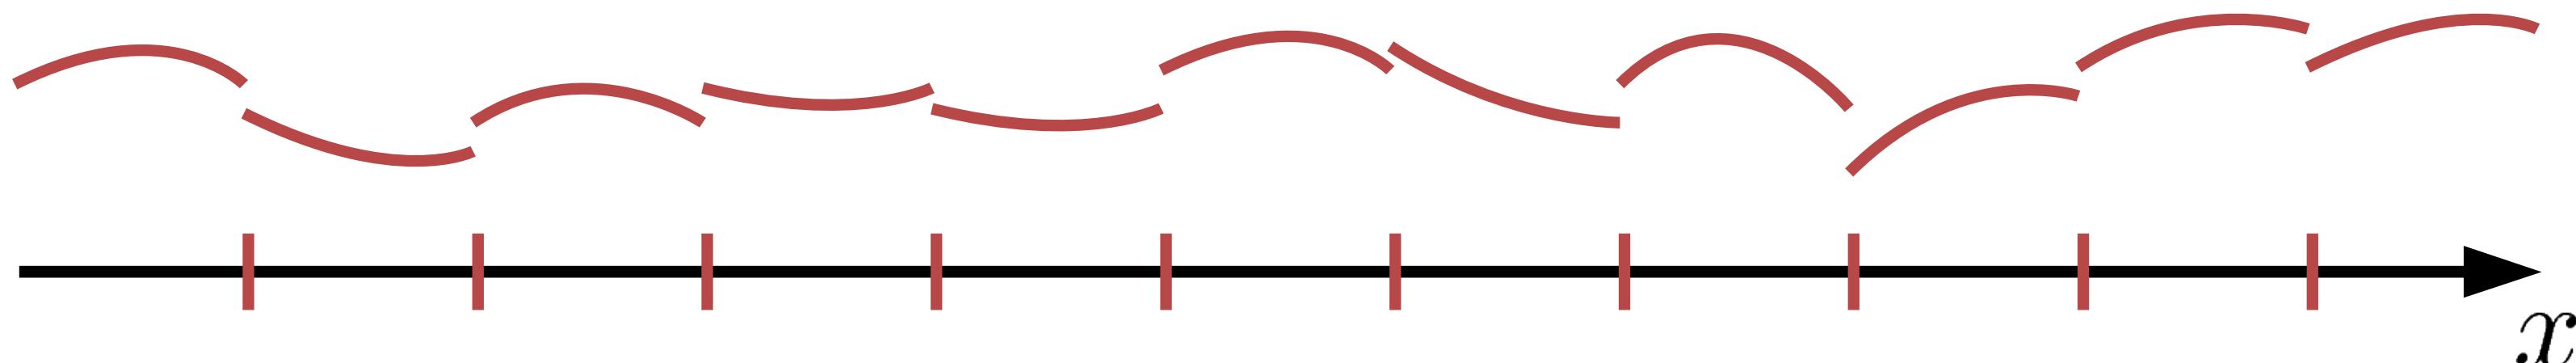
Theory

- Consider 1D scalar conservation law

$$\frac{\partial u}{\partial t} + \frac{\partial f}{\partial x} = 0$$

- Represent solution by order k piecewise discontinuous polynomials in each element

$$u^\delta = \bigoplus_{n=0}^{N-1} u_n^\delta \approx u$$



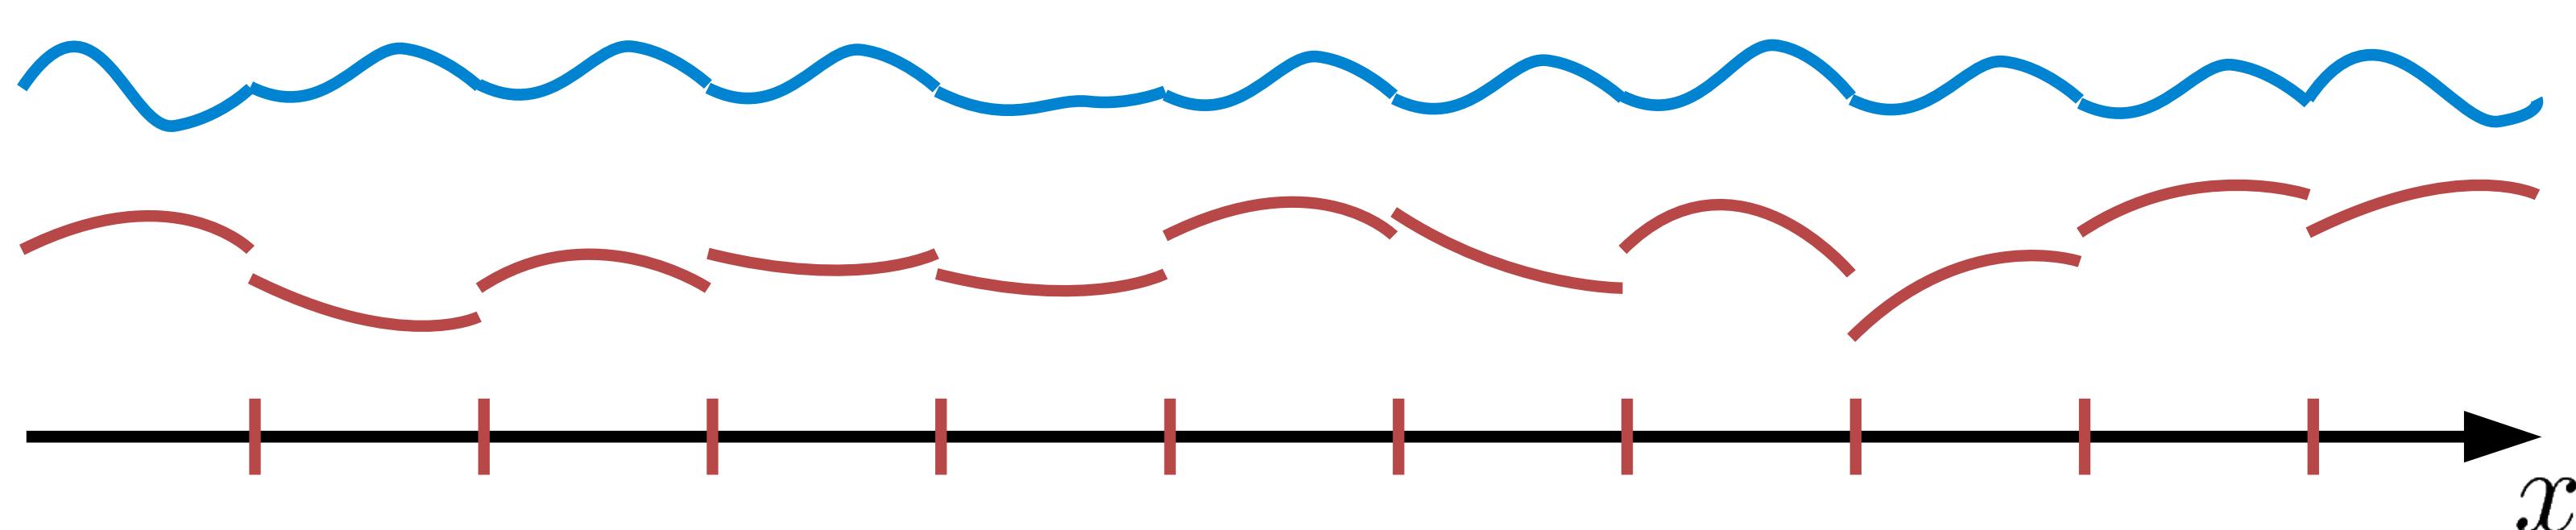
Theory

- Consider 1D scalar conservation law

$$\frac{\partial u}{\partial t} + \frac{\partial f}{\partial x} = 0$$

- Represent flux by order $k+1$ piecewise continuous polynomials within each element

$$f^\delta = \bigoplus_{n=0}^{N-1} f_n^\delta \approx f$$



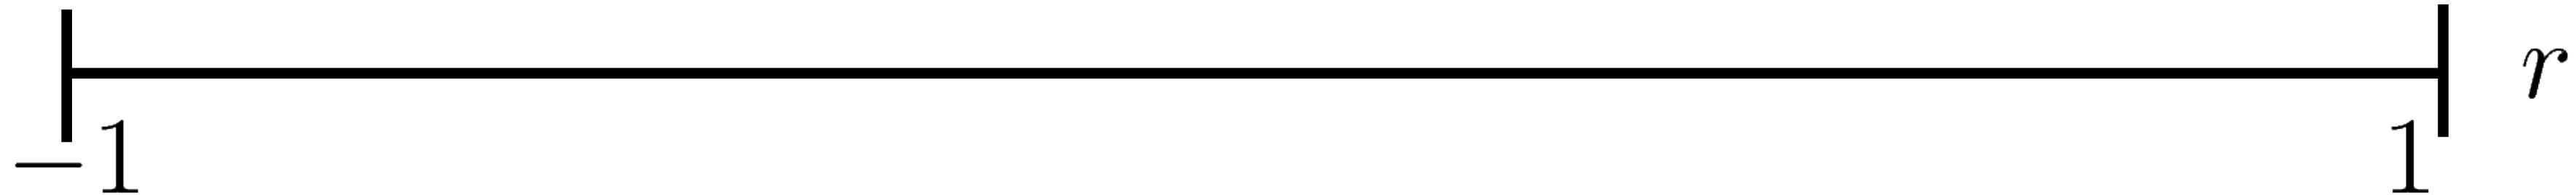
Theory

- Map each element to a standard element

$$\Omega_S = \{r \mid -1 \leq r \leq 1\} \quad r = 2 \left(\frac{x - x_n}{x_{n+1} - x_n} \right) - 1$$

$$\frac{\partial \hat{u}^\delta}{\partial t} + \frac{\partial \hat{f}^\delta}{\partial r} = 0 \quad \hat{u}^\delta = u_n^\delta \quad \hat{f}^\delta = \frac{f_n^\delta}{h_n}$$

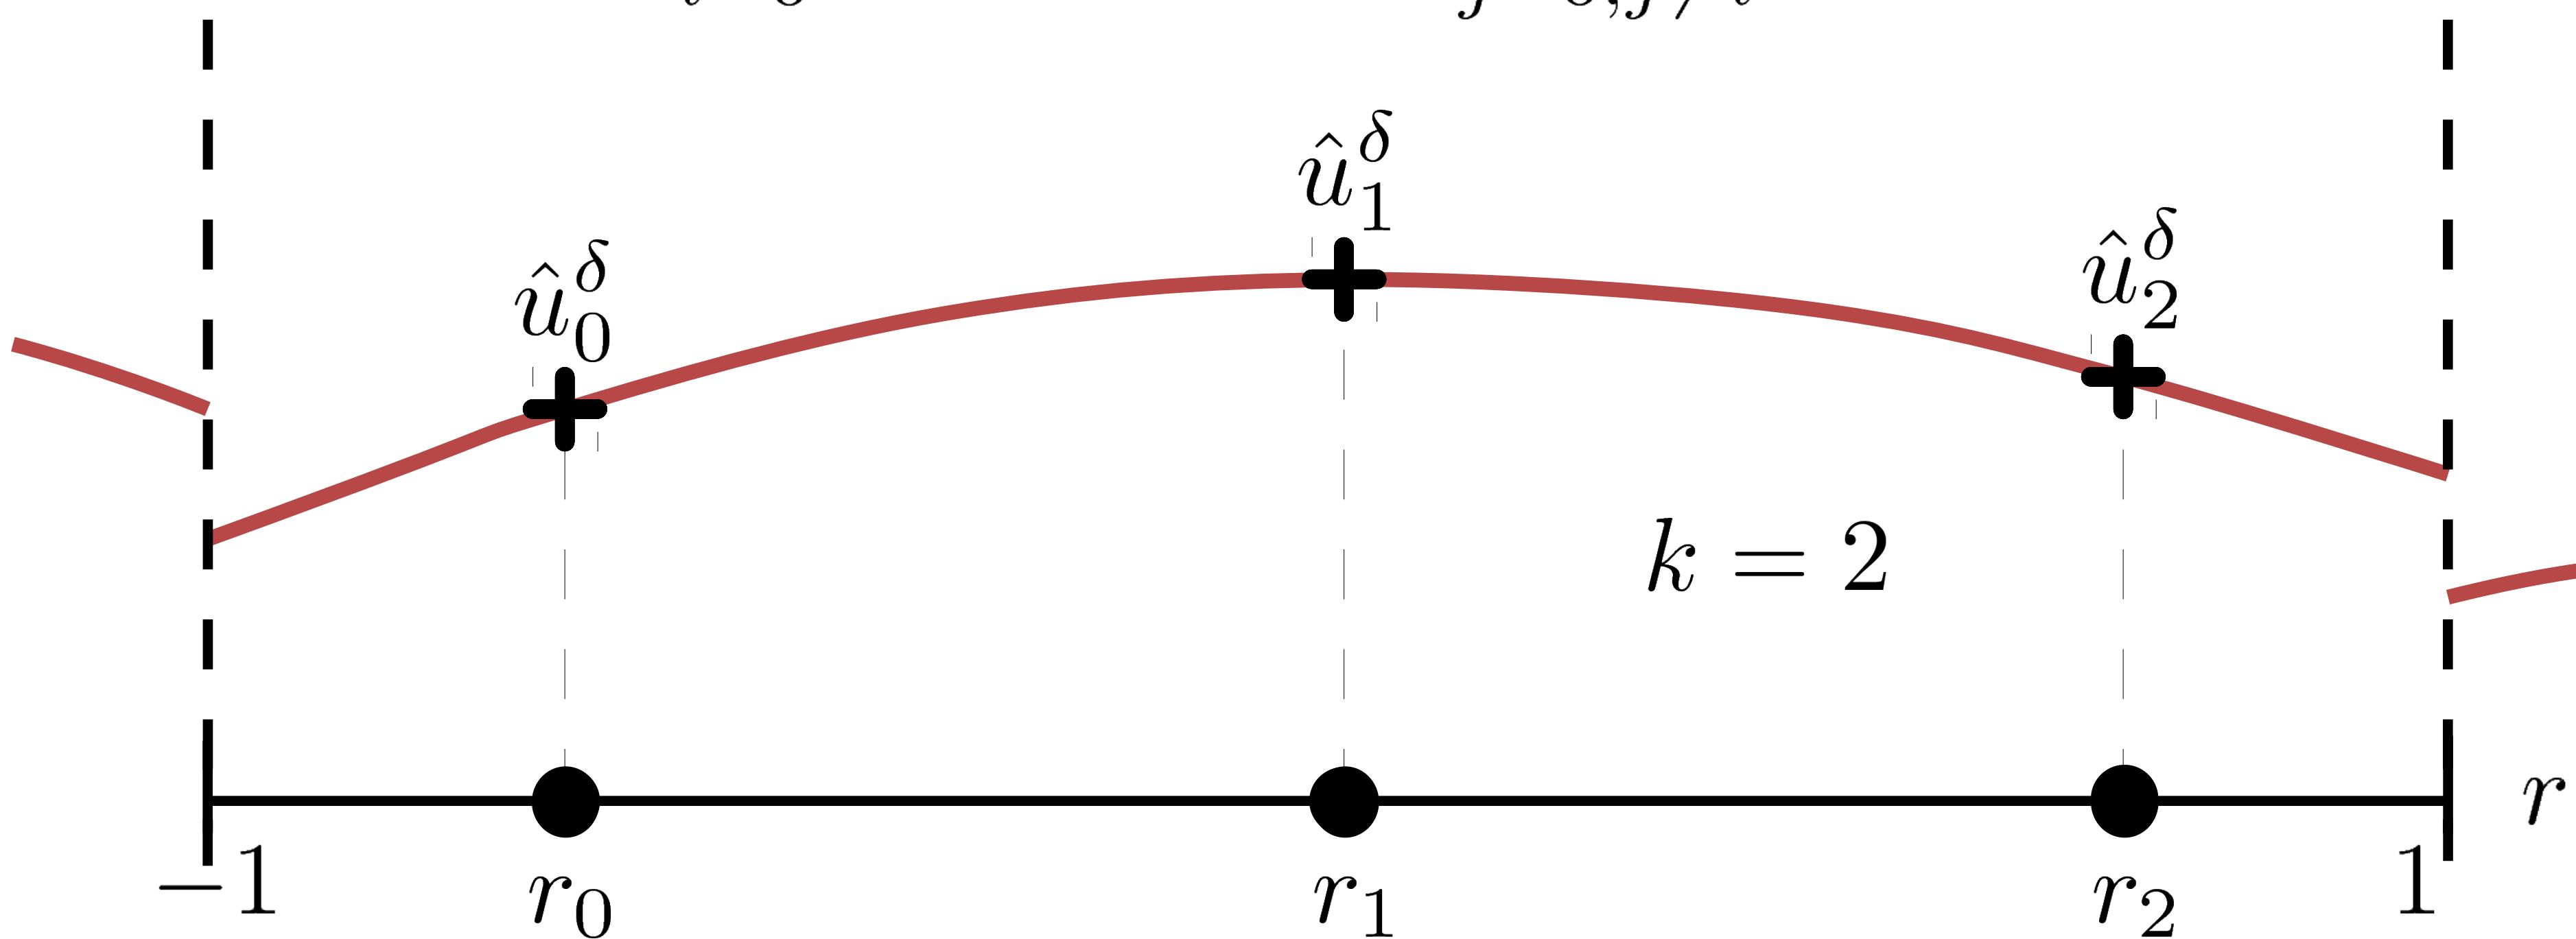
$$h_n = (x_{n+1} - x_n)/2$$



Theory

- Represent order k solution within standard element using a nodal basis

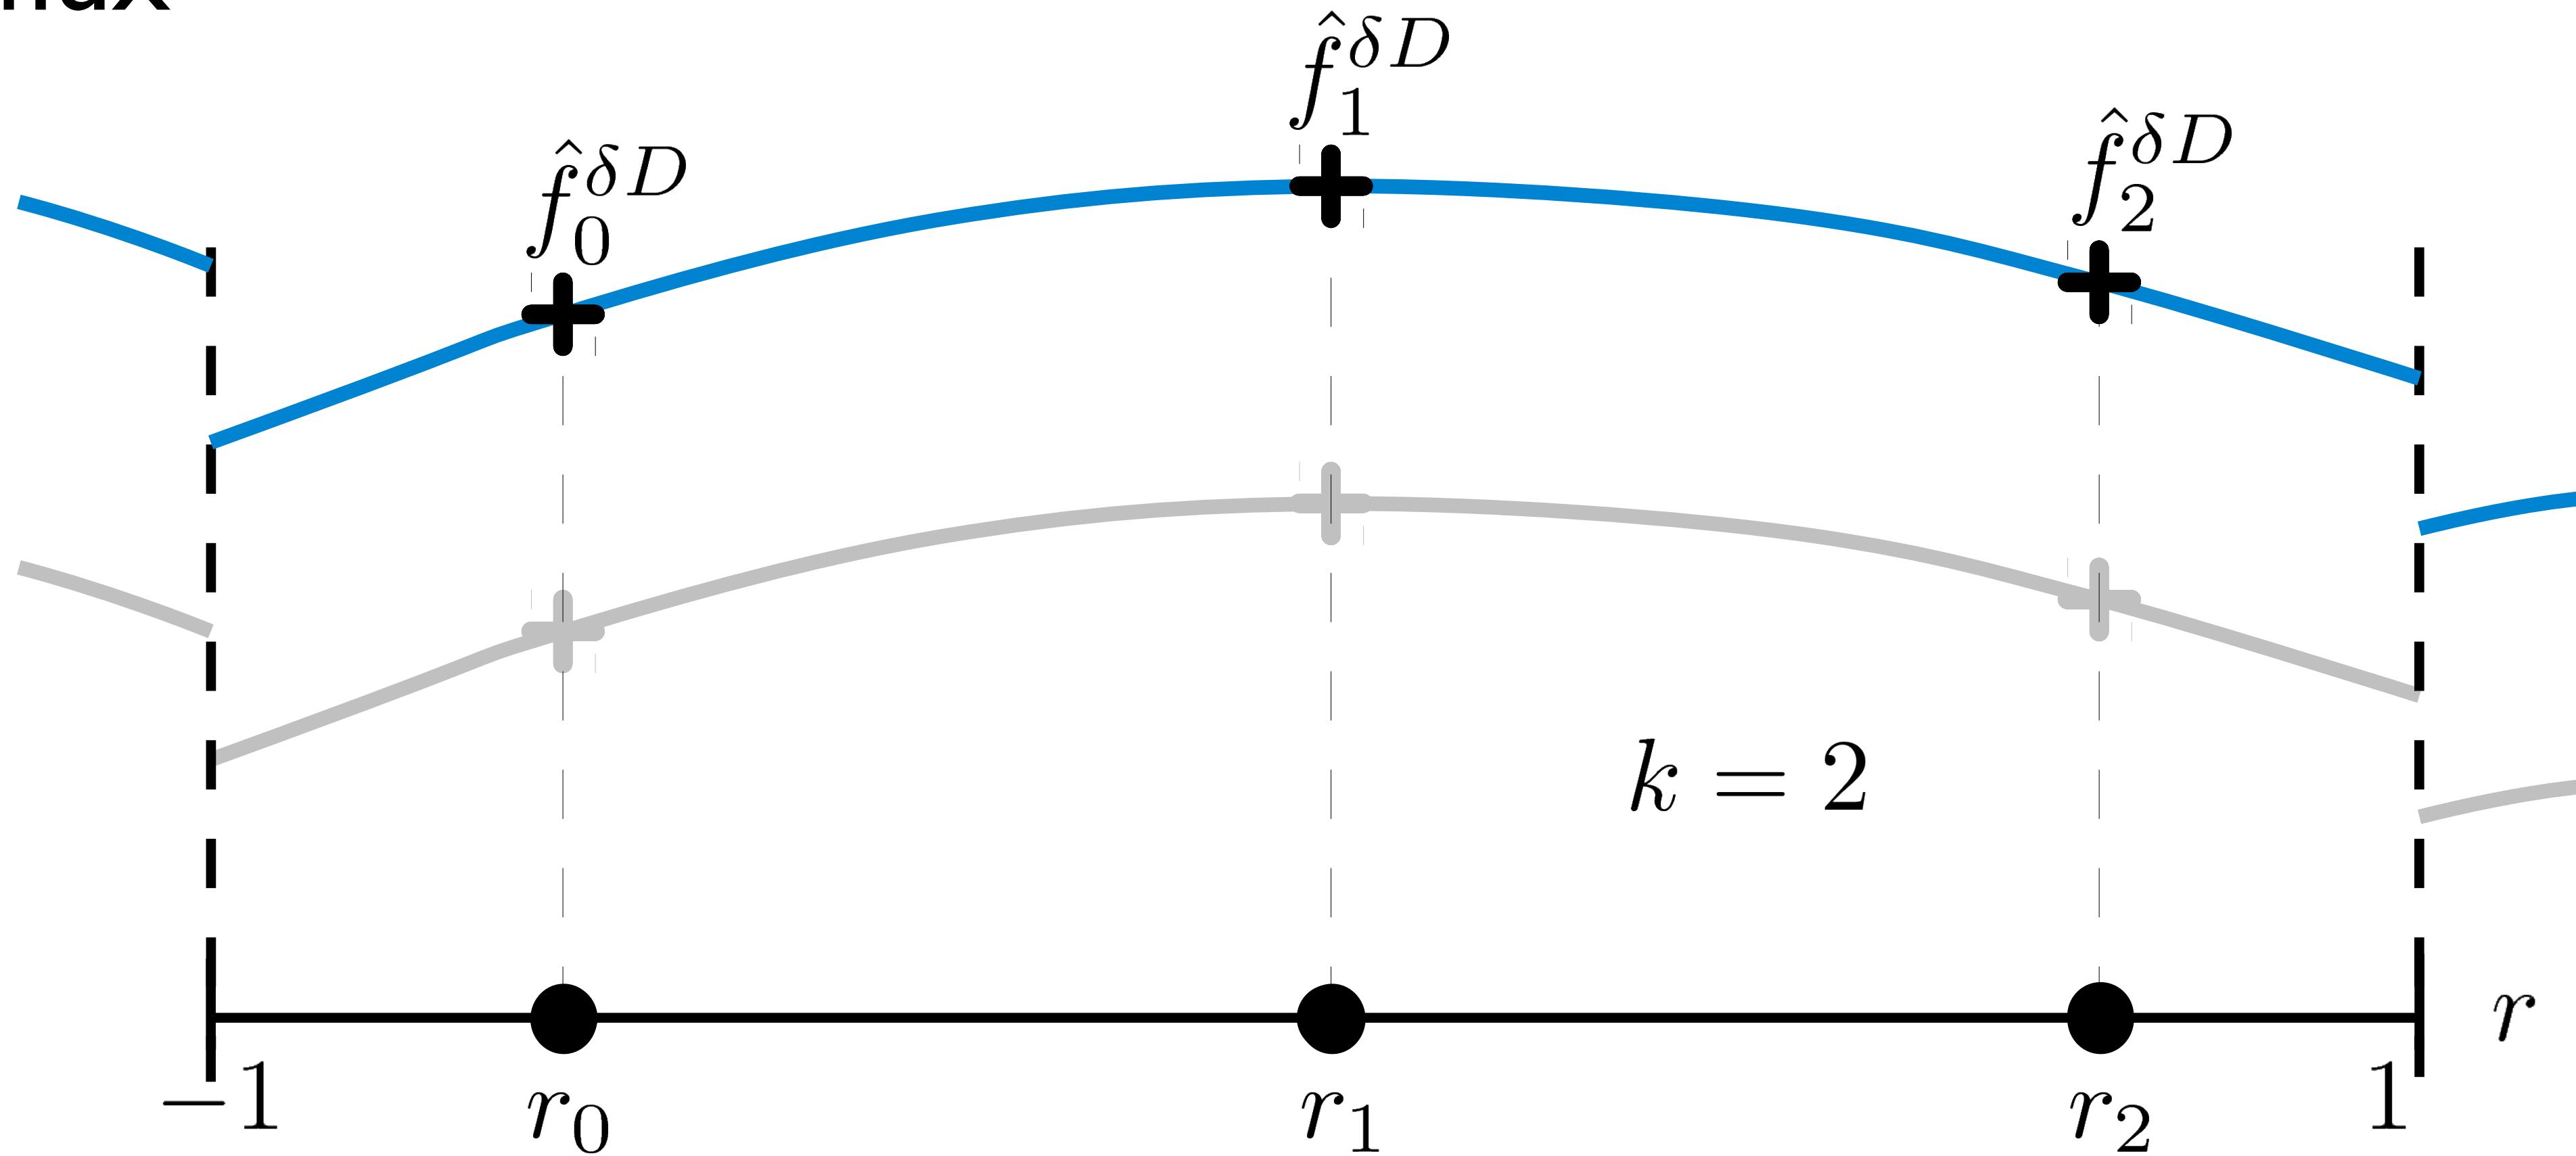
$$\hat{u}^\delta = \sum_{i=0}^k \hat{u}_i^\delta l_i \quad l_i = \prod_{j=0, j \neq i}^k \left(\frac{r - r_j}{r_i - r_j} \right)$$



Theory

- Reconstruct discontinuous flux

$$\hat{f}^{\delta D} = \sum_{i=0}^k \hat{f}_i^{\delta D} l_i$$

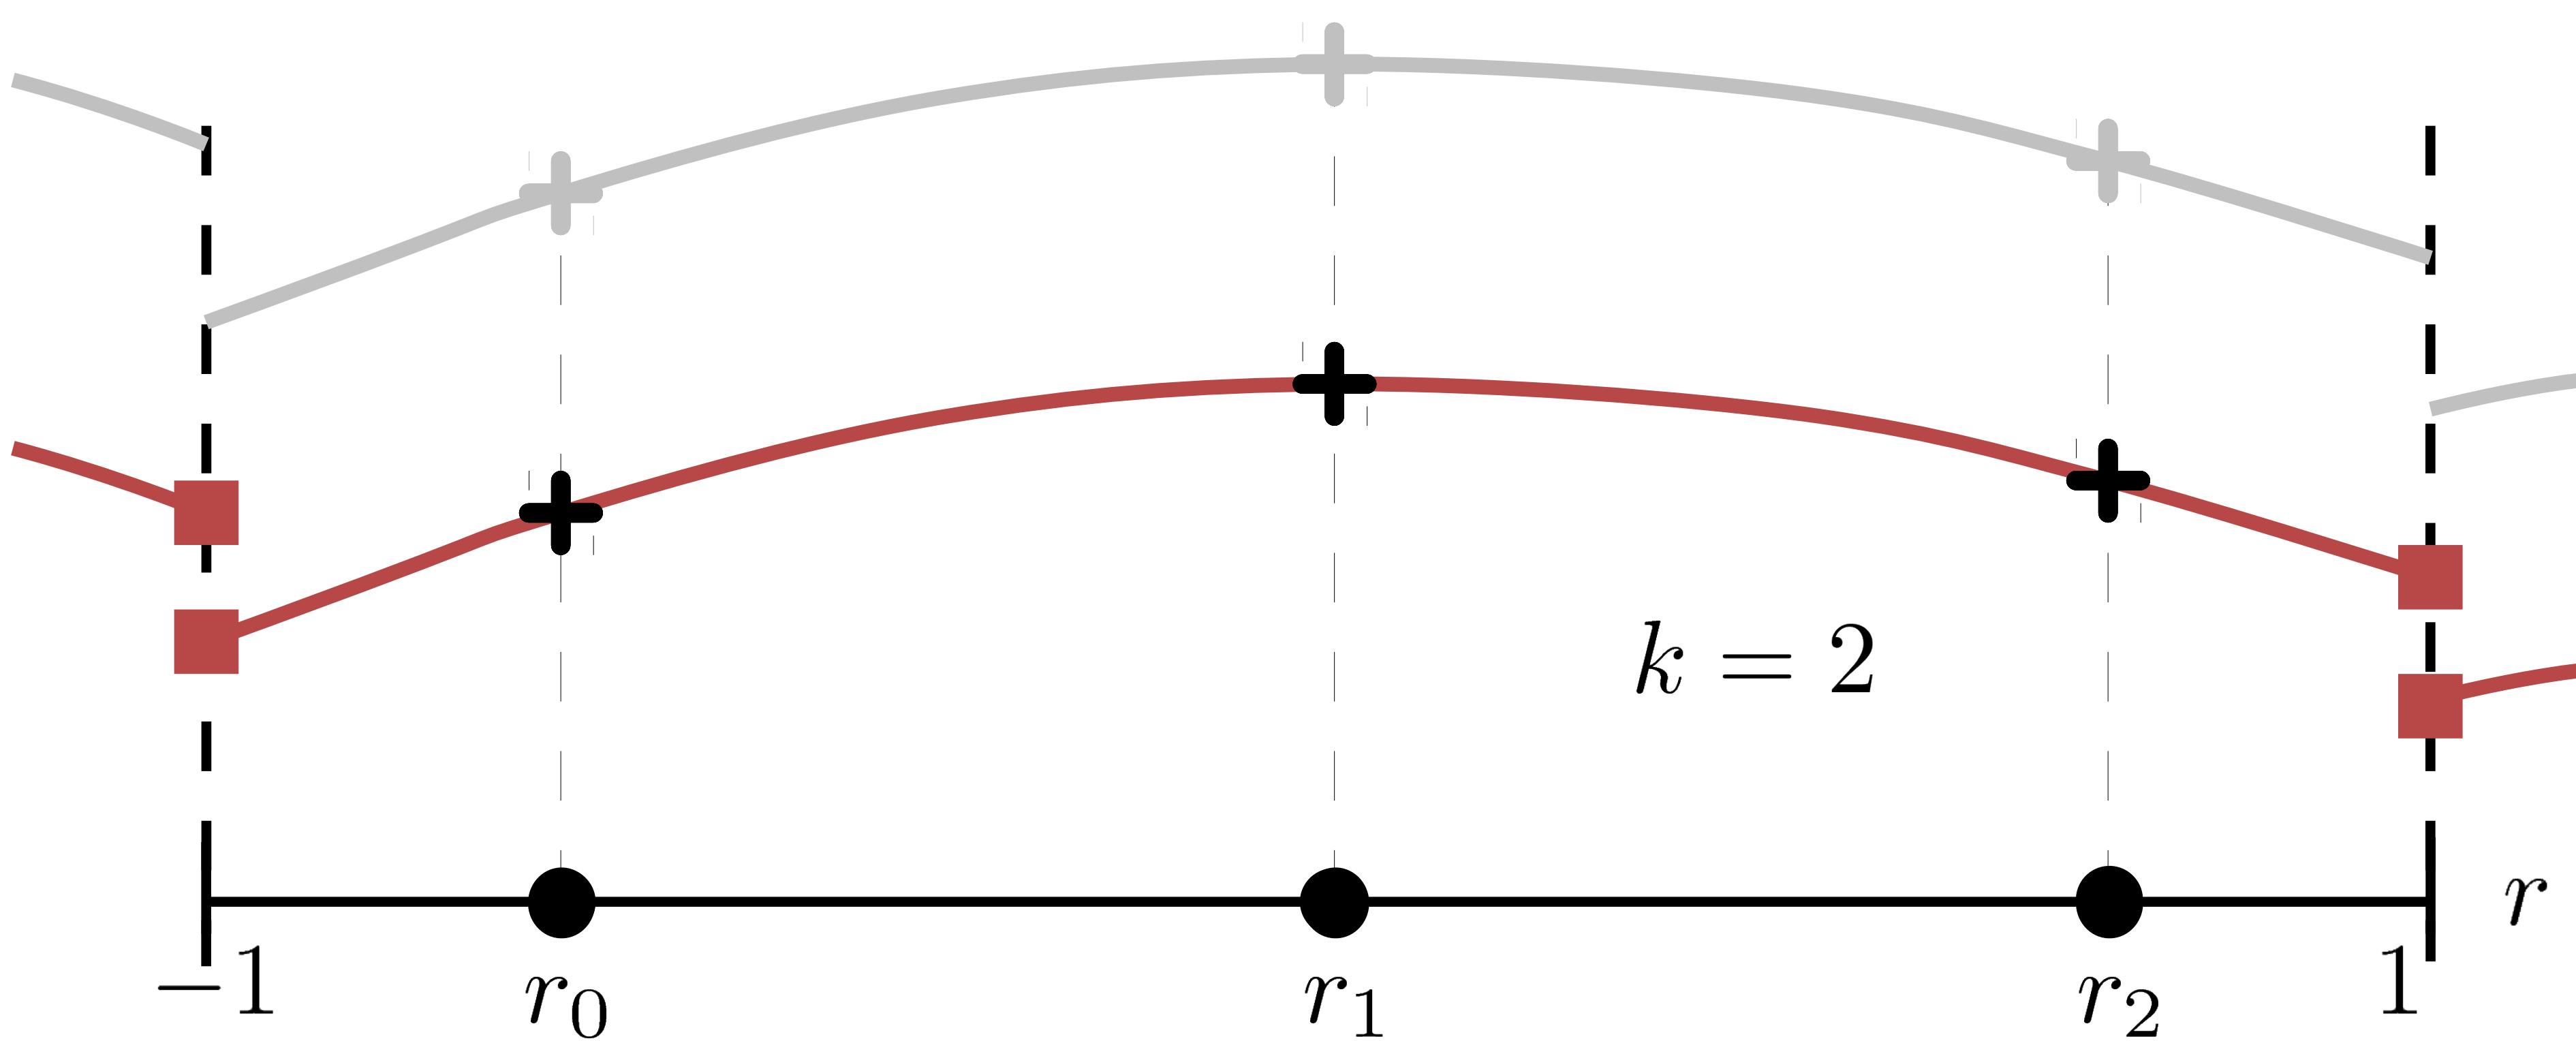


Theory

- Evaluate solution at element boundaries

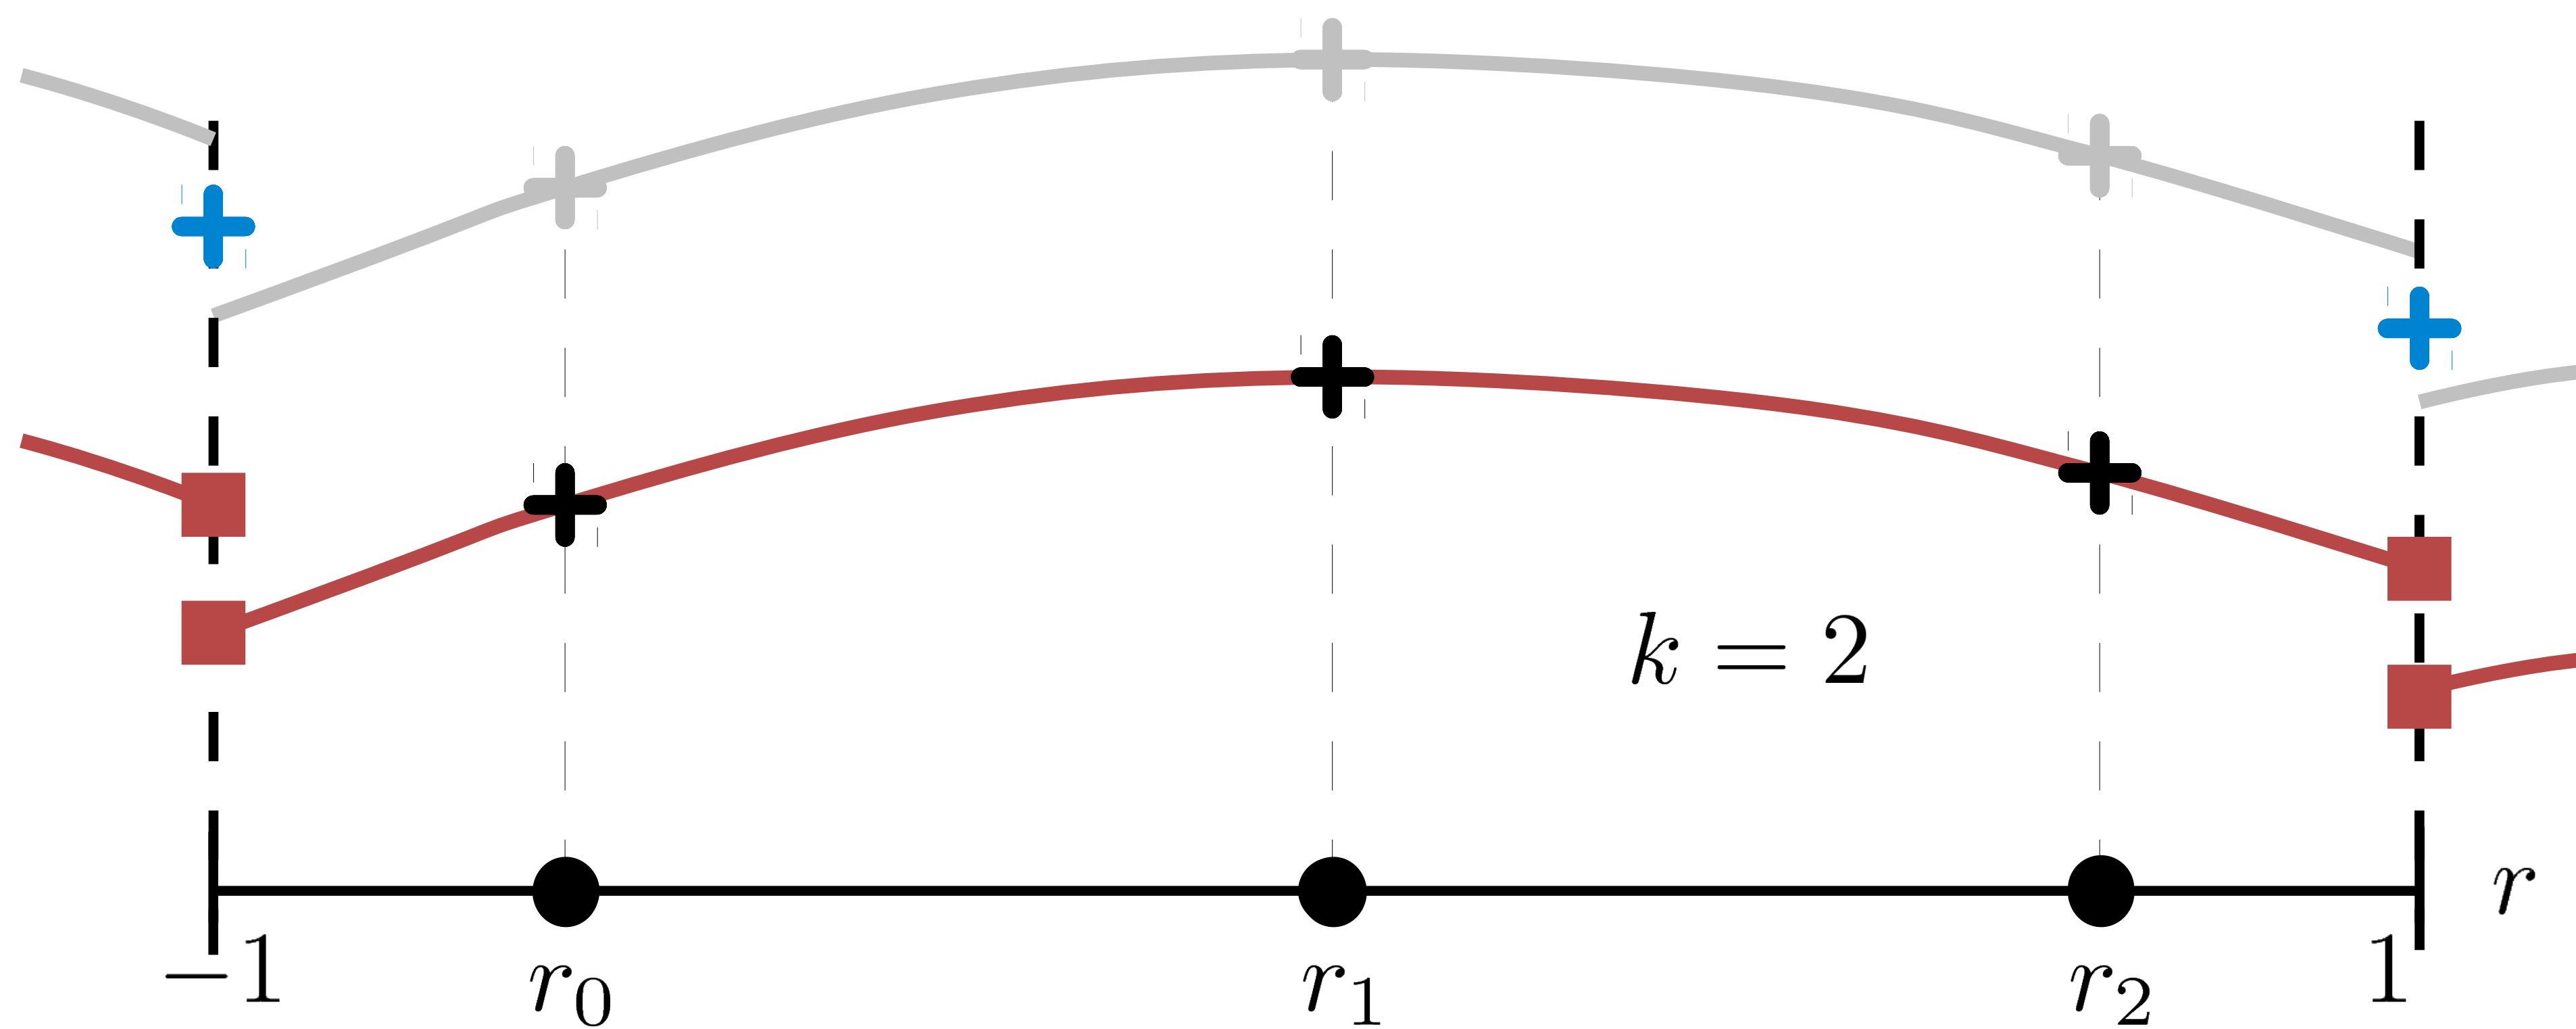
$$\hat{u}^\delta(-1) = \sum_{i=0}^k \hat{u}_i^\delta l_i(-1)$$

$$\hat{u}^\delta(1) = \sum_{i=0}^k \hat{u}_i^\delta l_i(1)$$



Theory

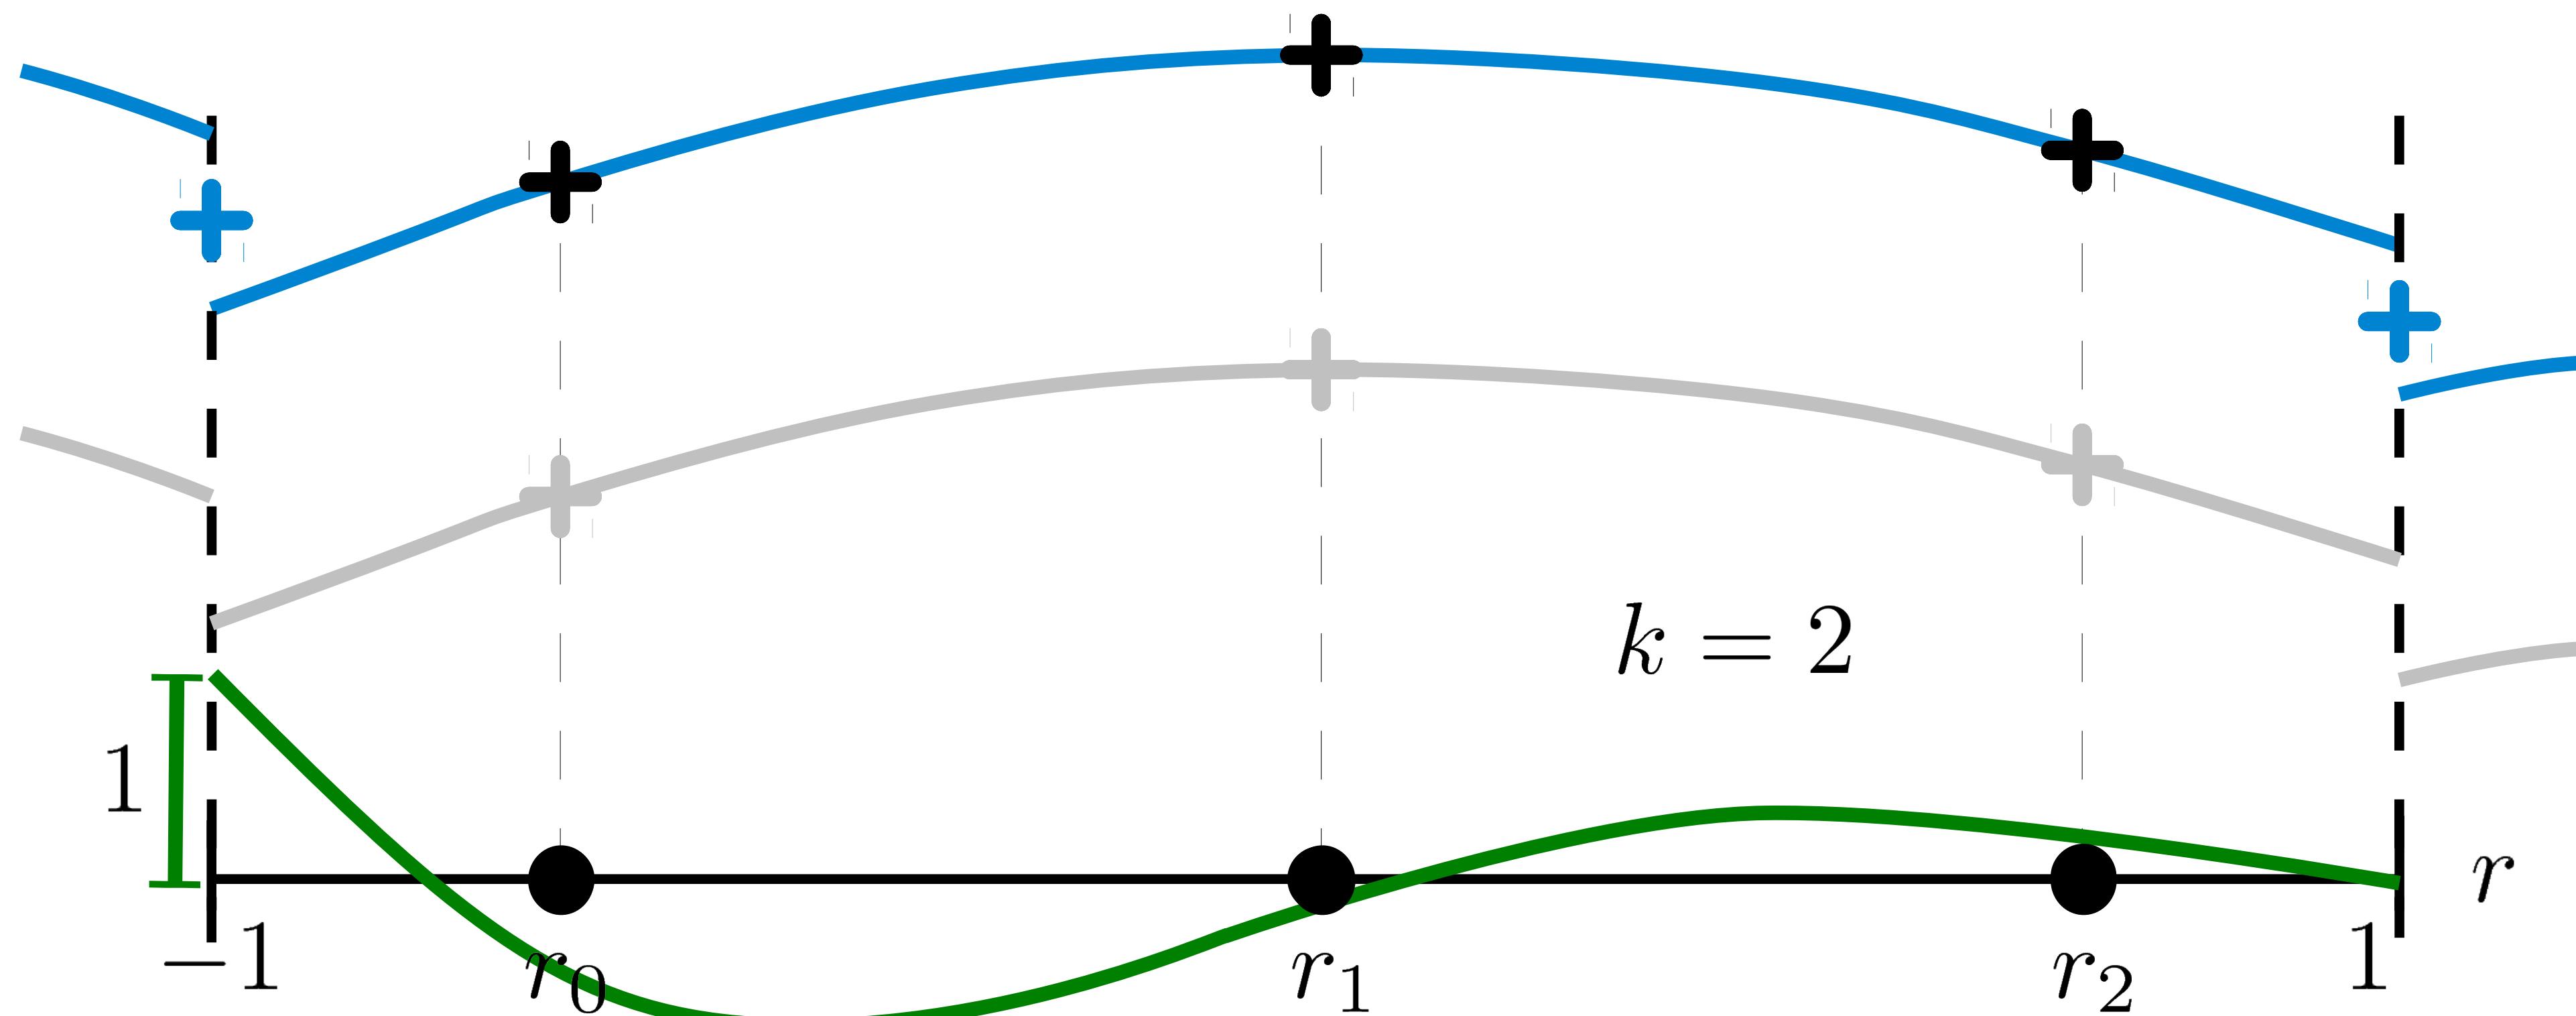
- Calculate common interface fluxes



Theory

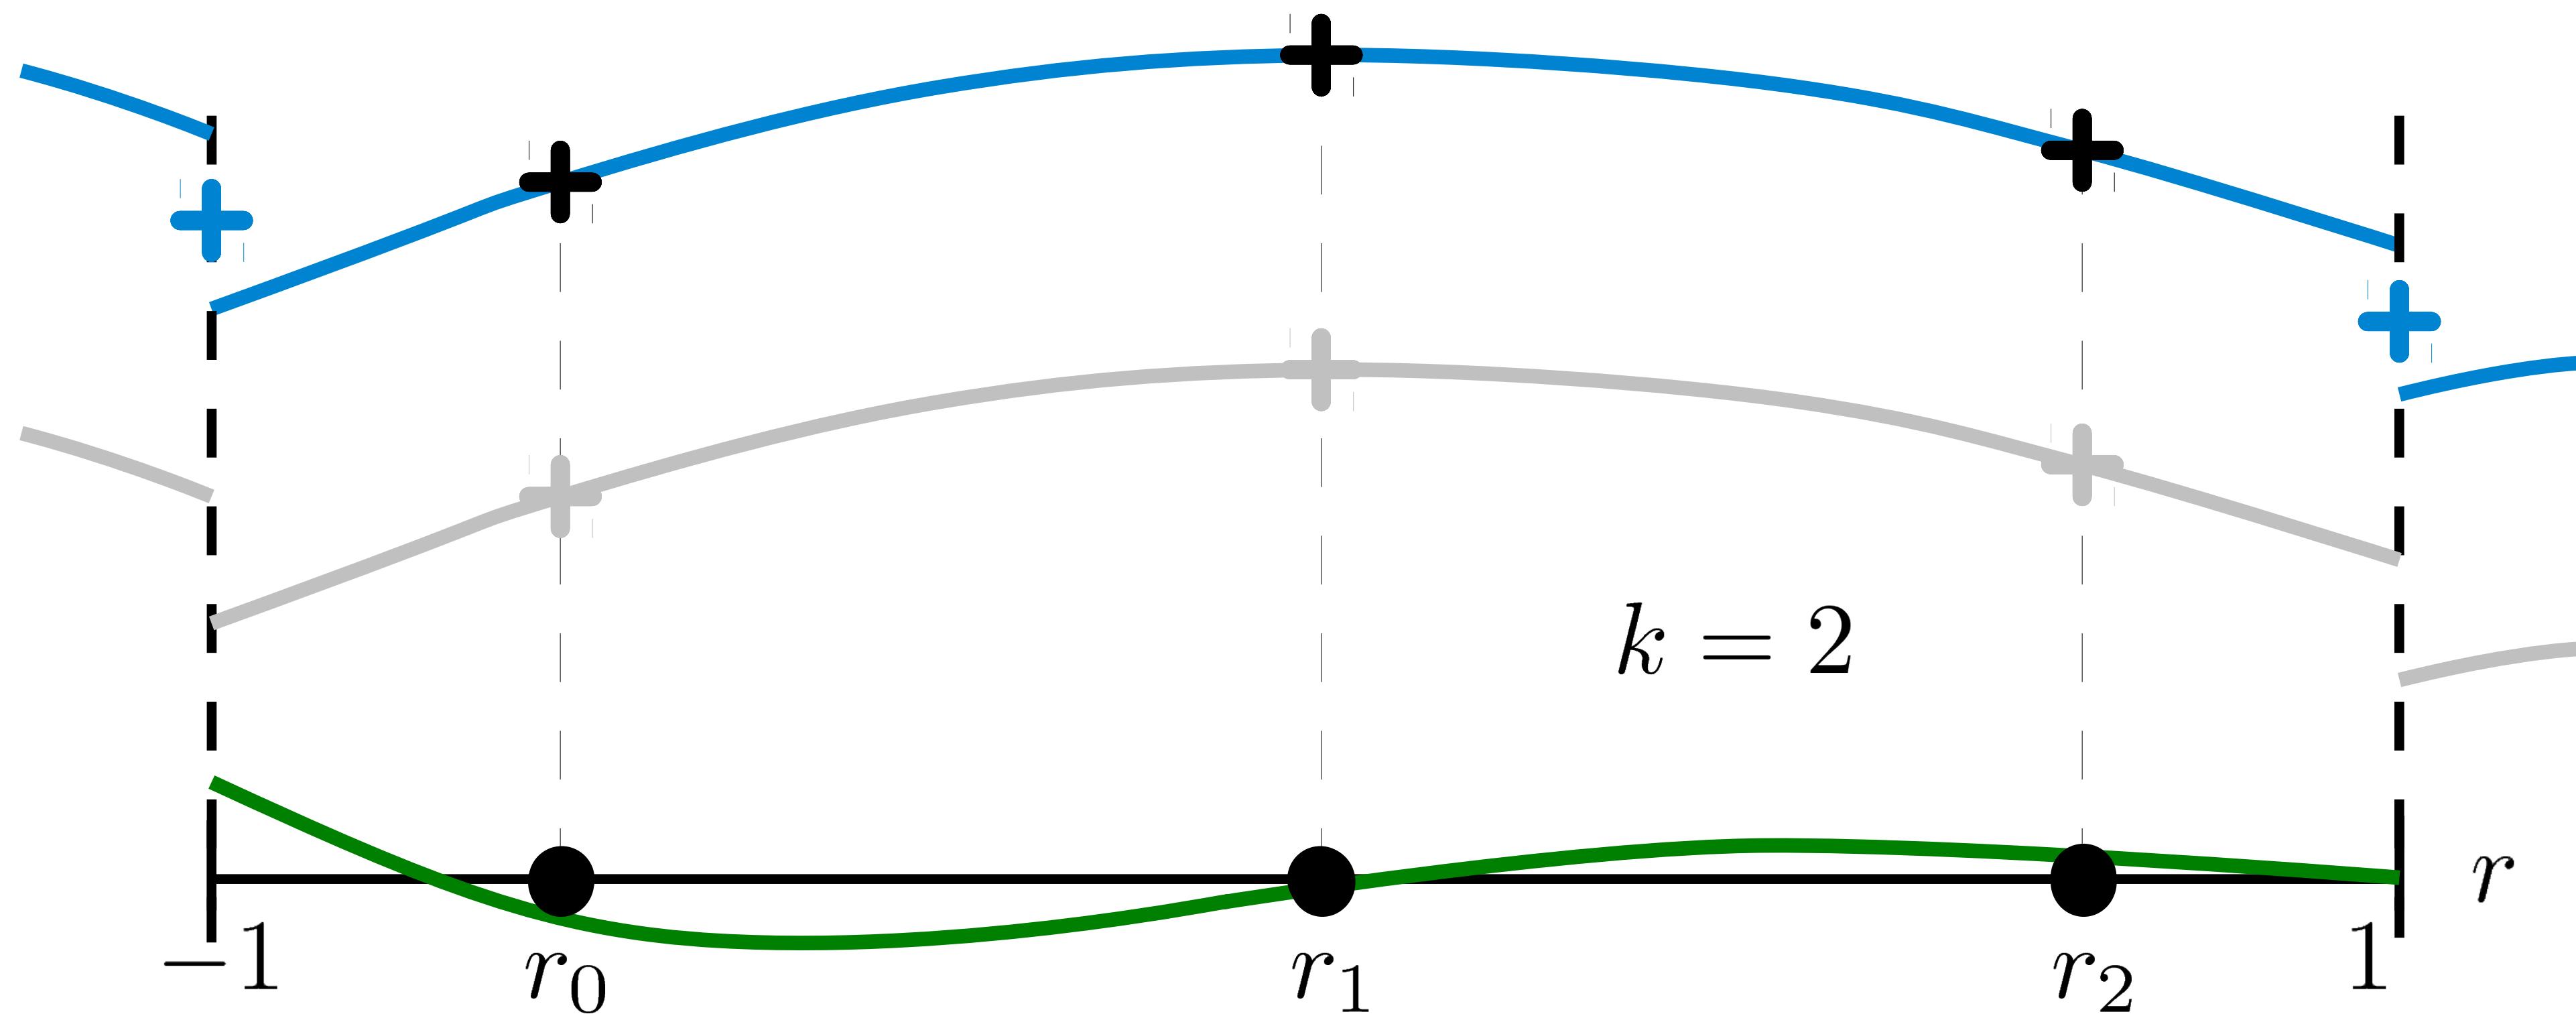
- Define order $k+1$ ‘correction function’

$$g_L(-1) = 1, \quad g_L(1) = 0$$



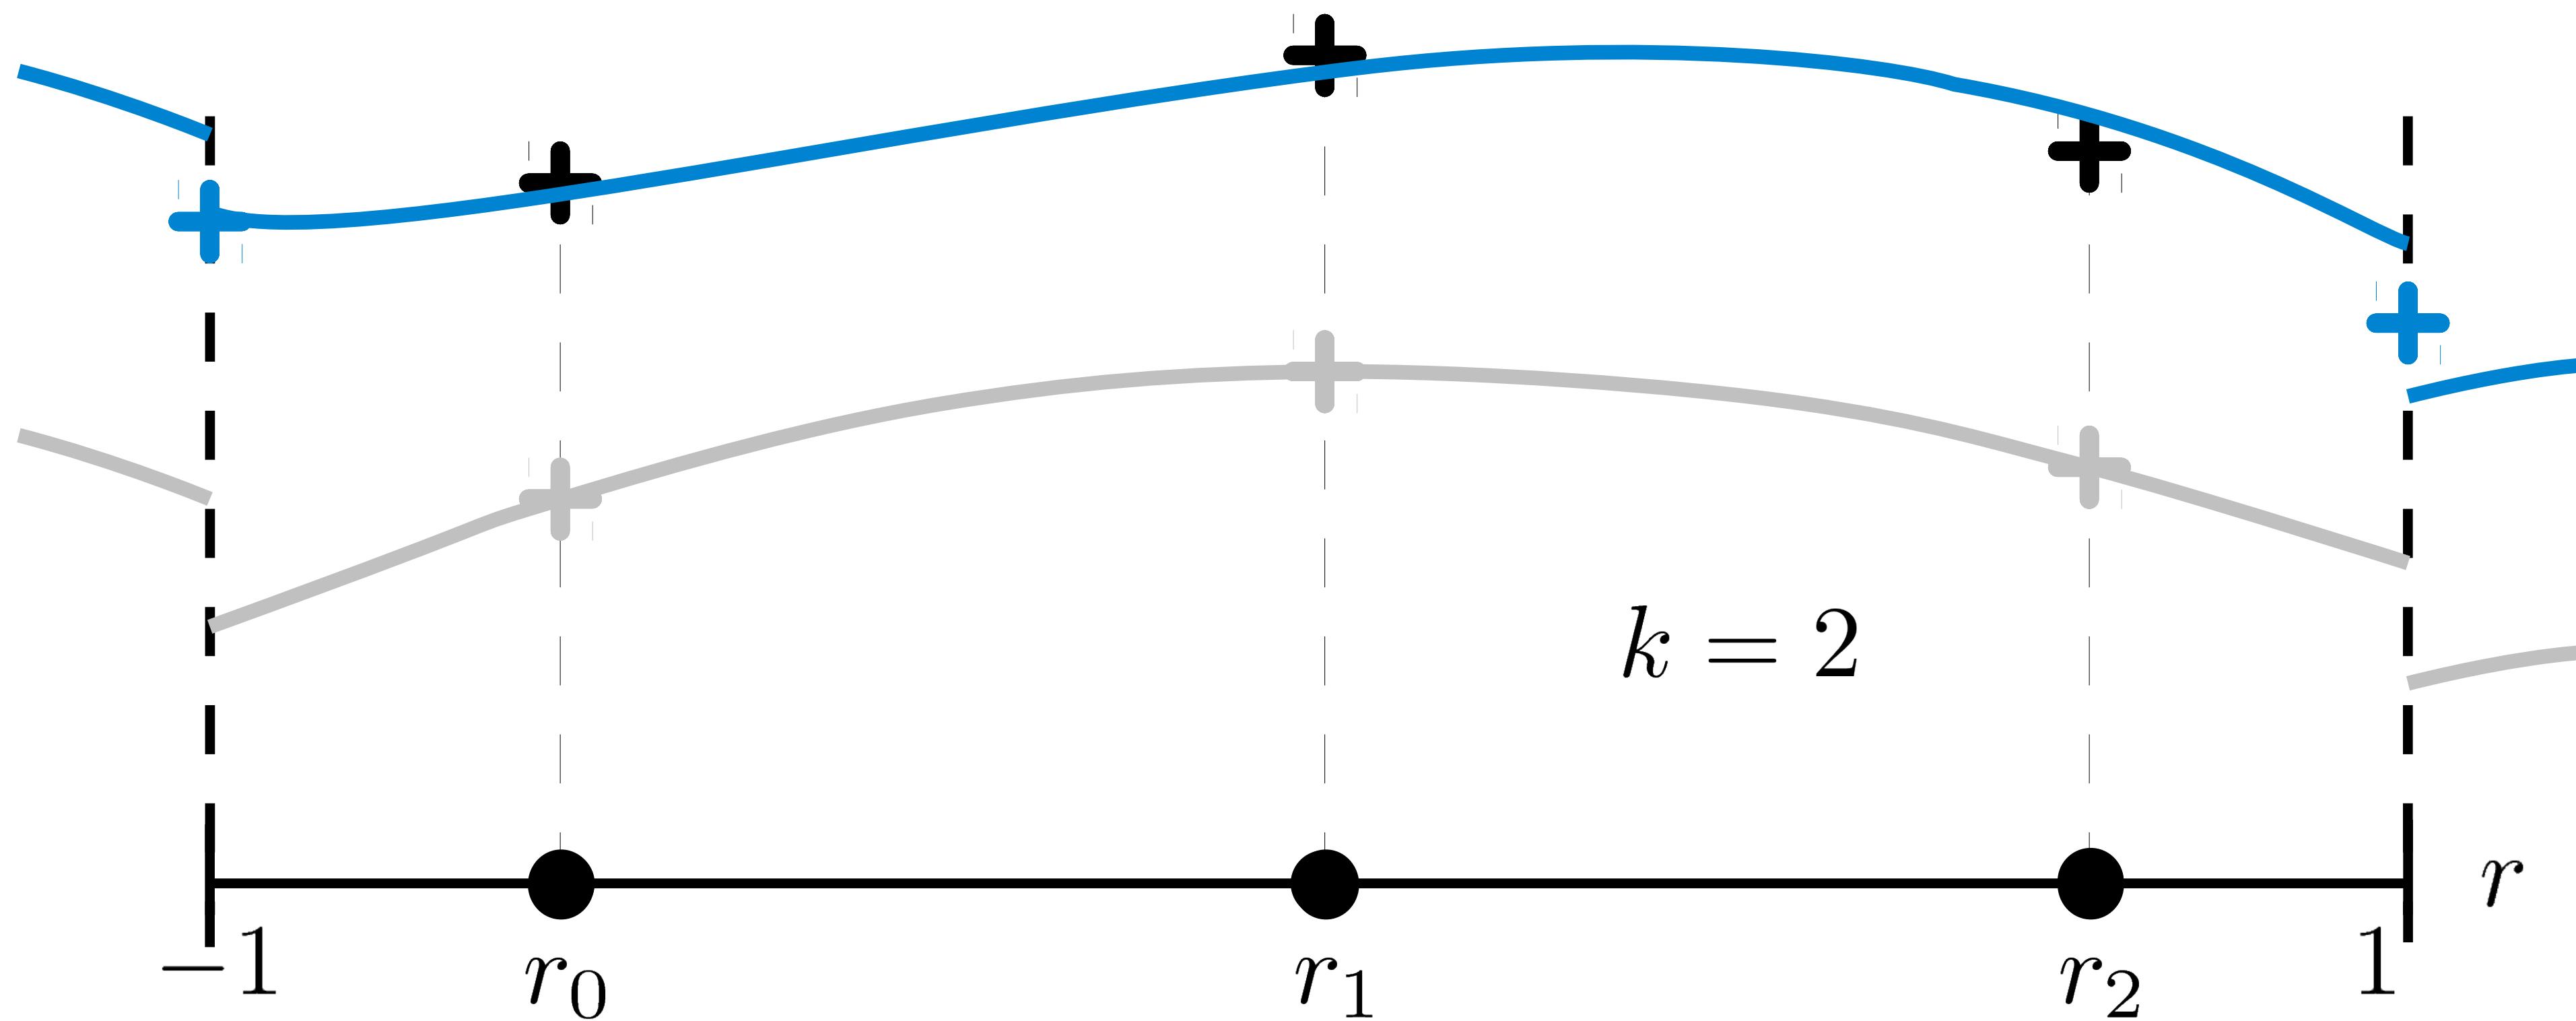
Theory

- Scale this correction function...



Theory

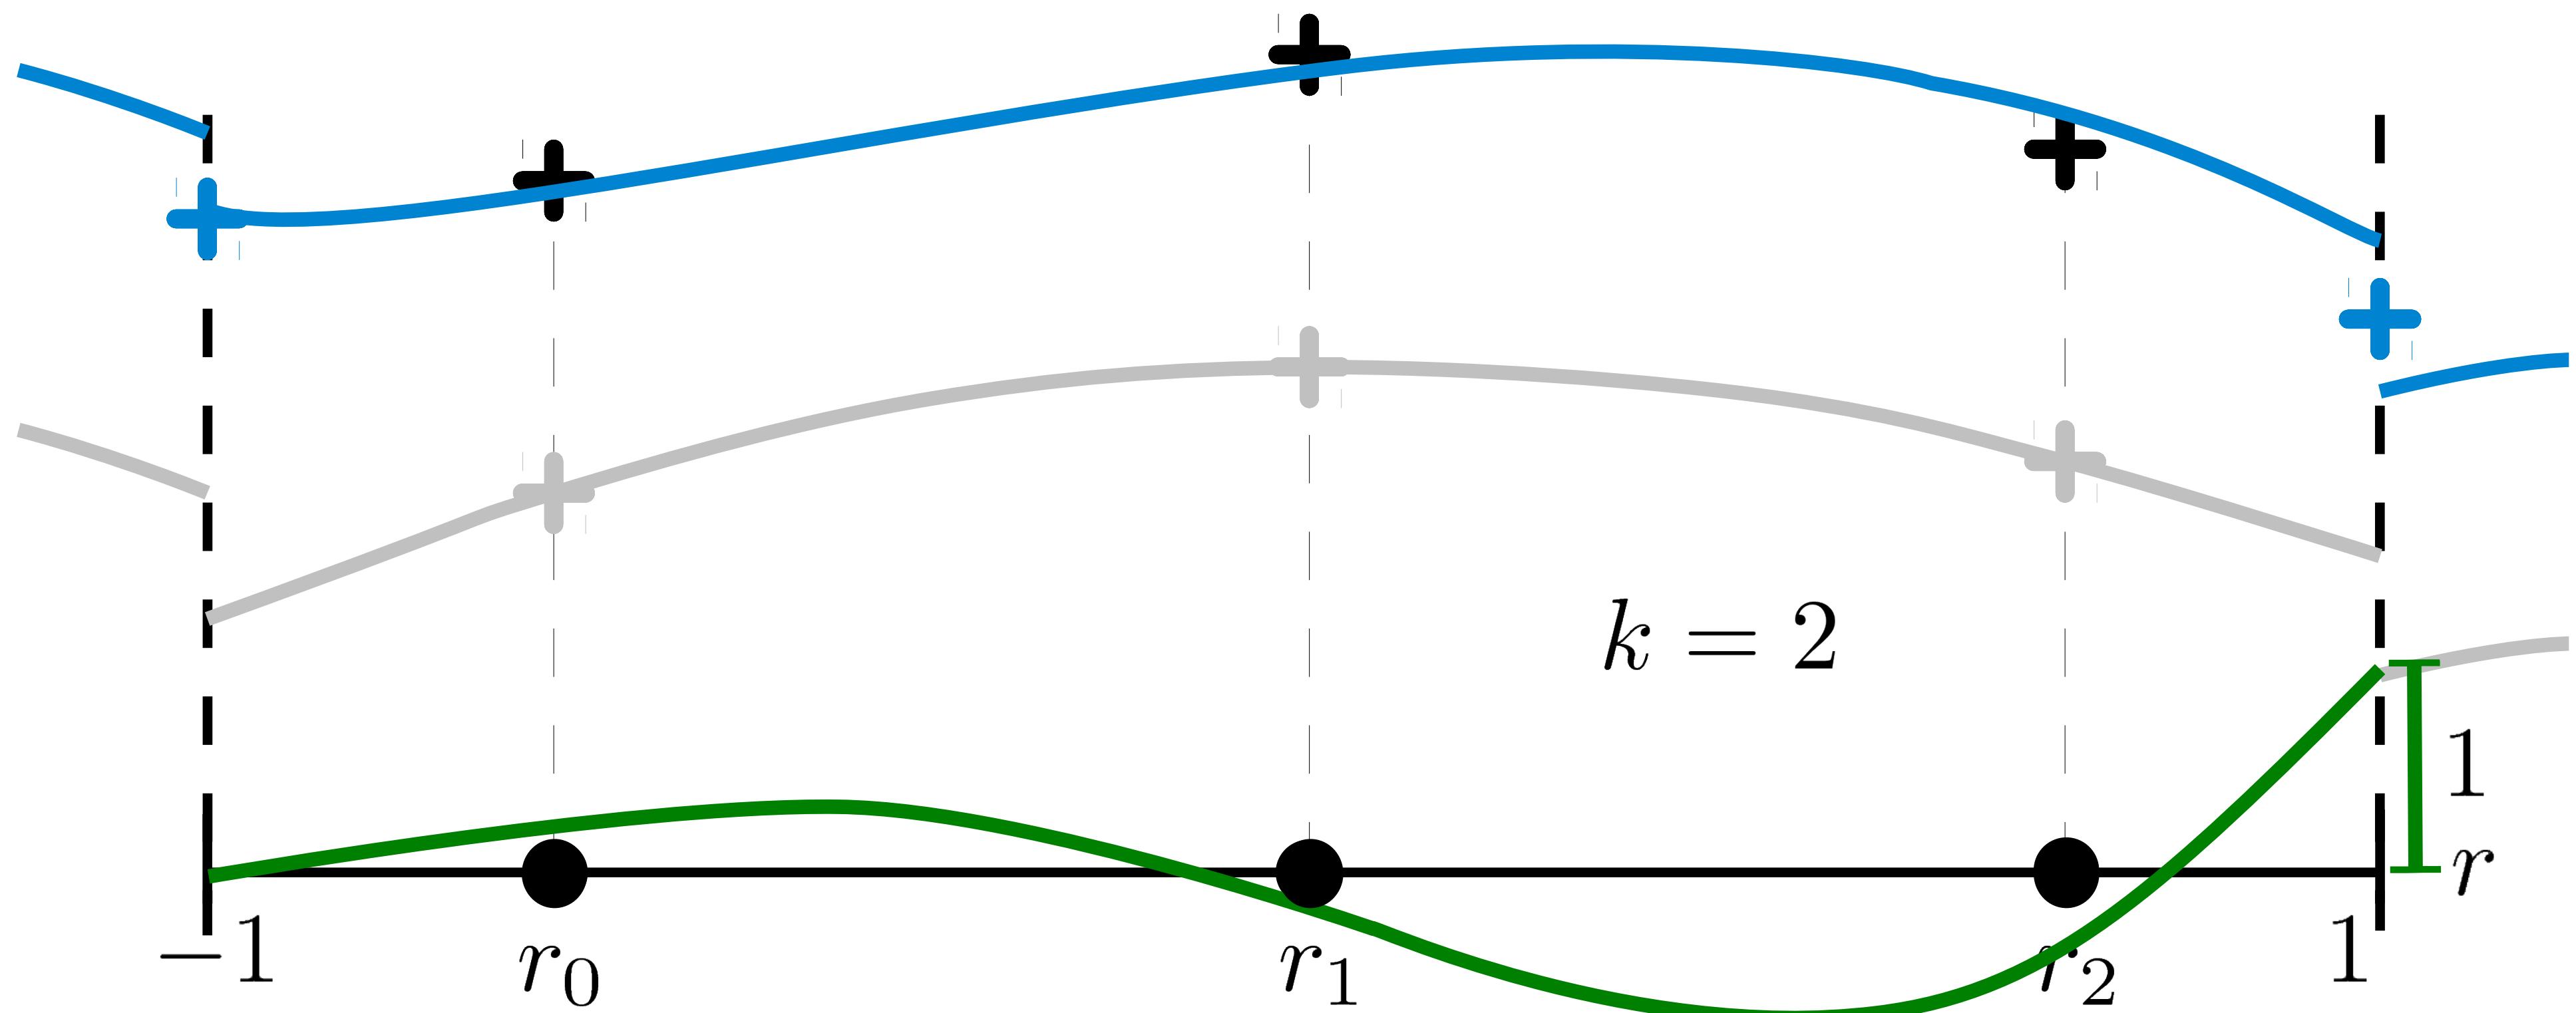
- ... and add it to the discontinuous flux



Theory

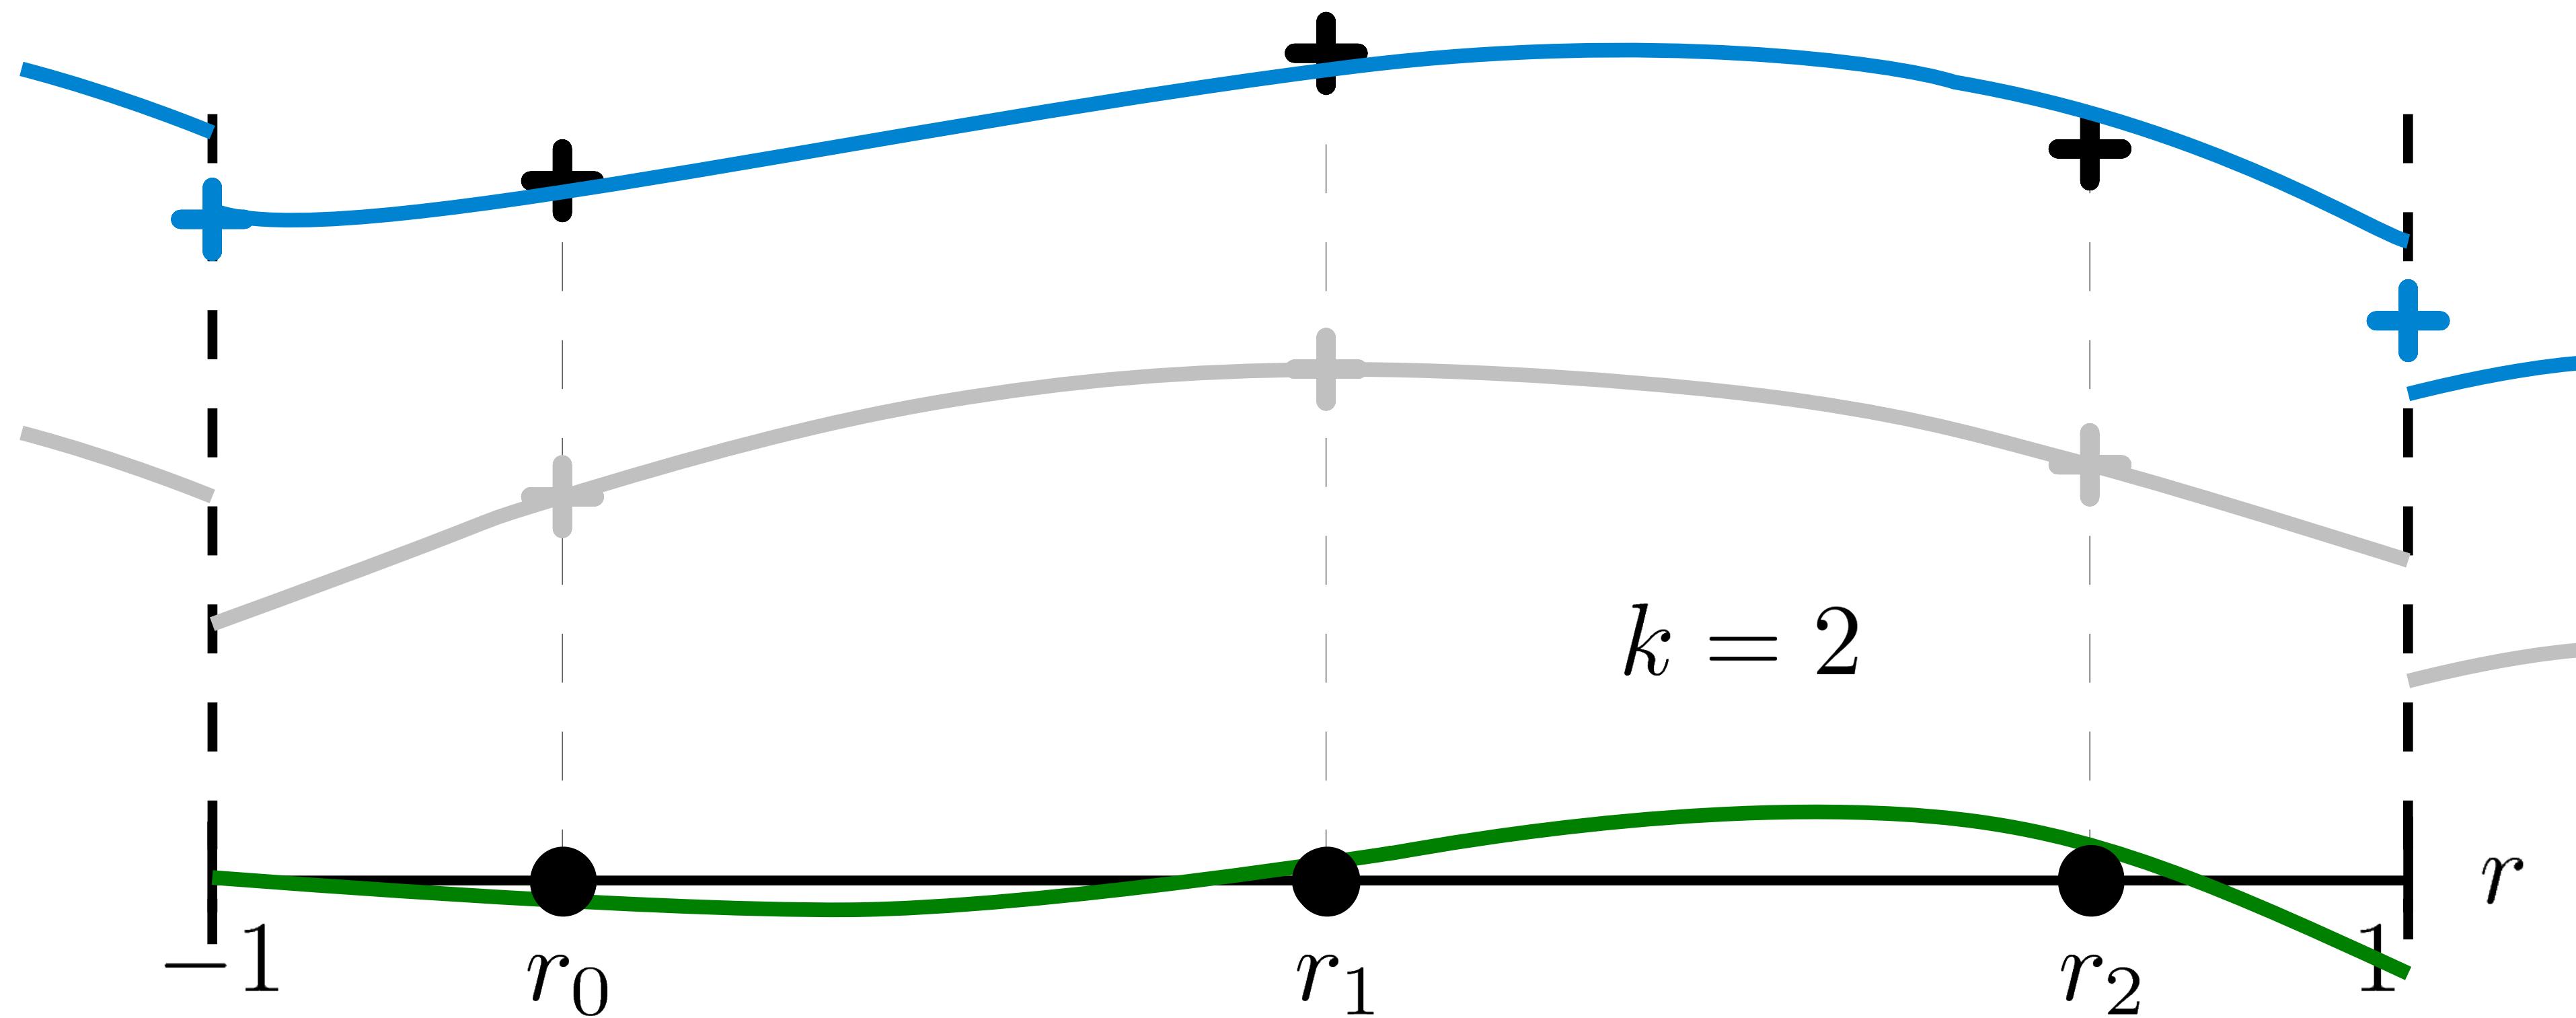
- Repeat from the right...

$$g_R(r) = g_L(-r)$$



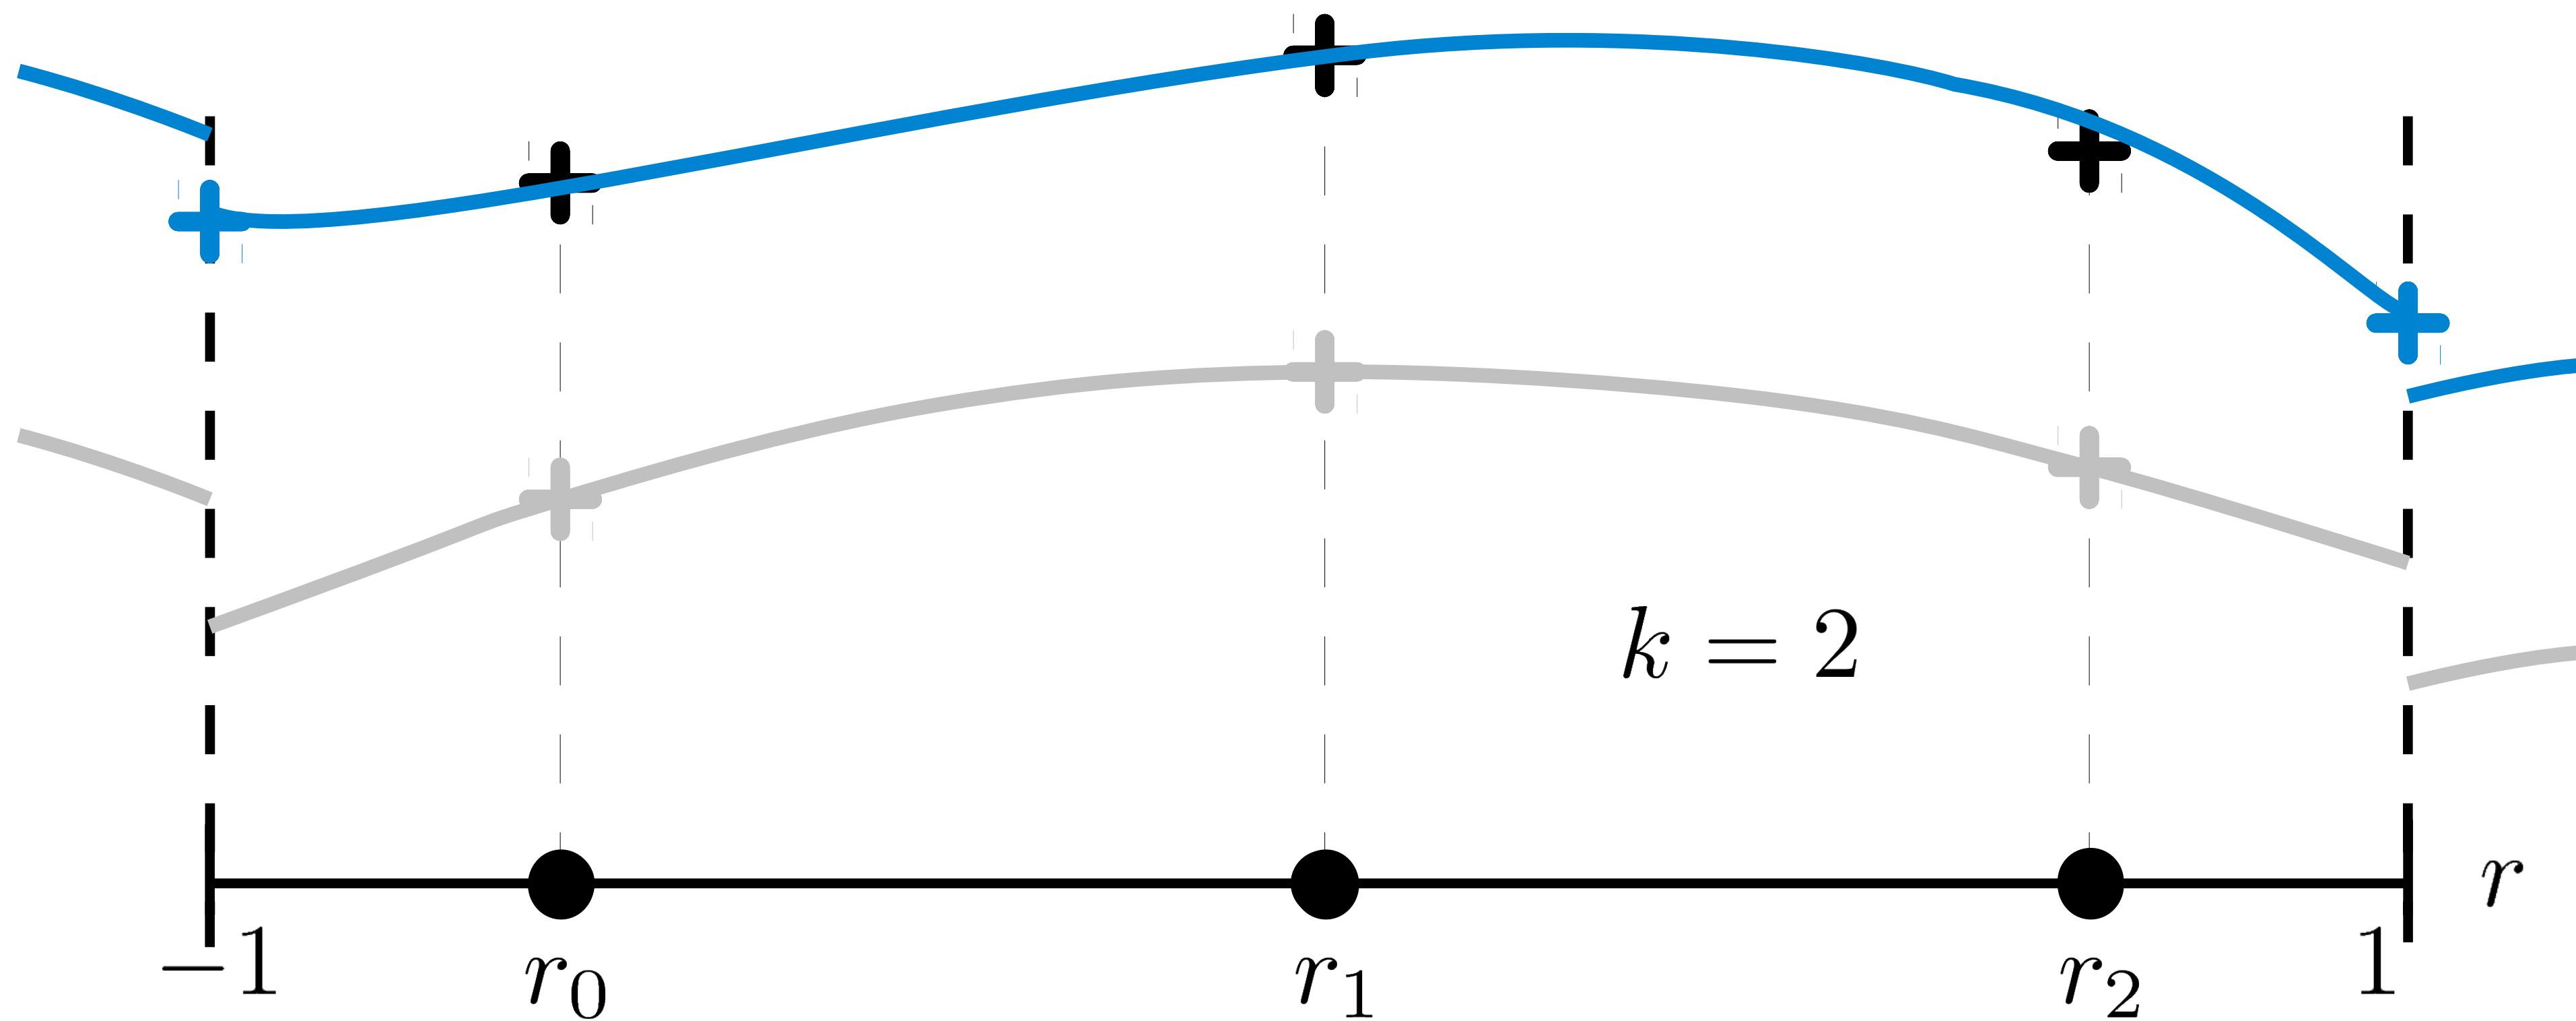
Theory

- Repeat from the right...



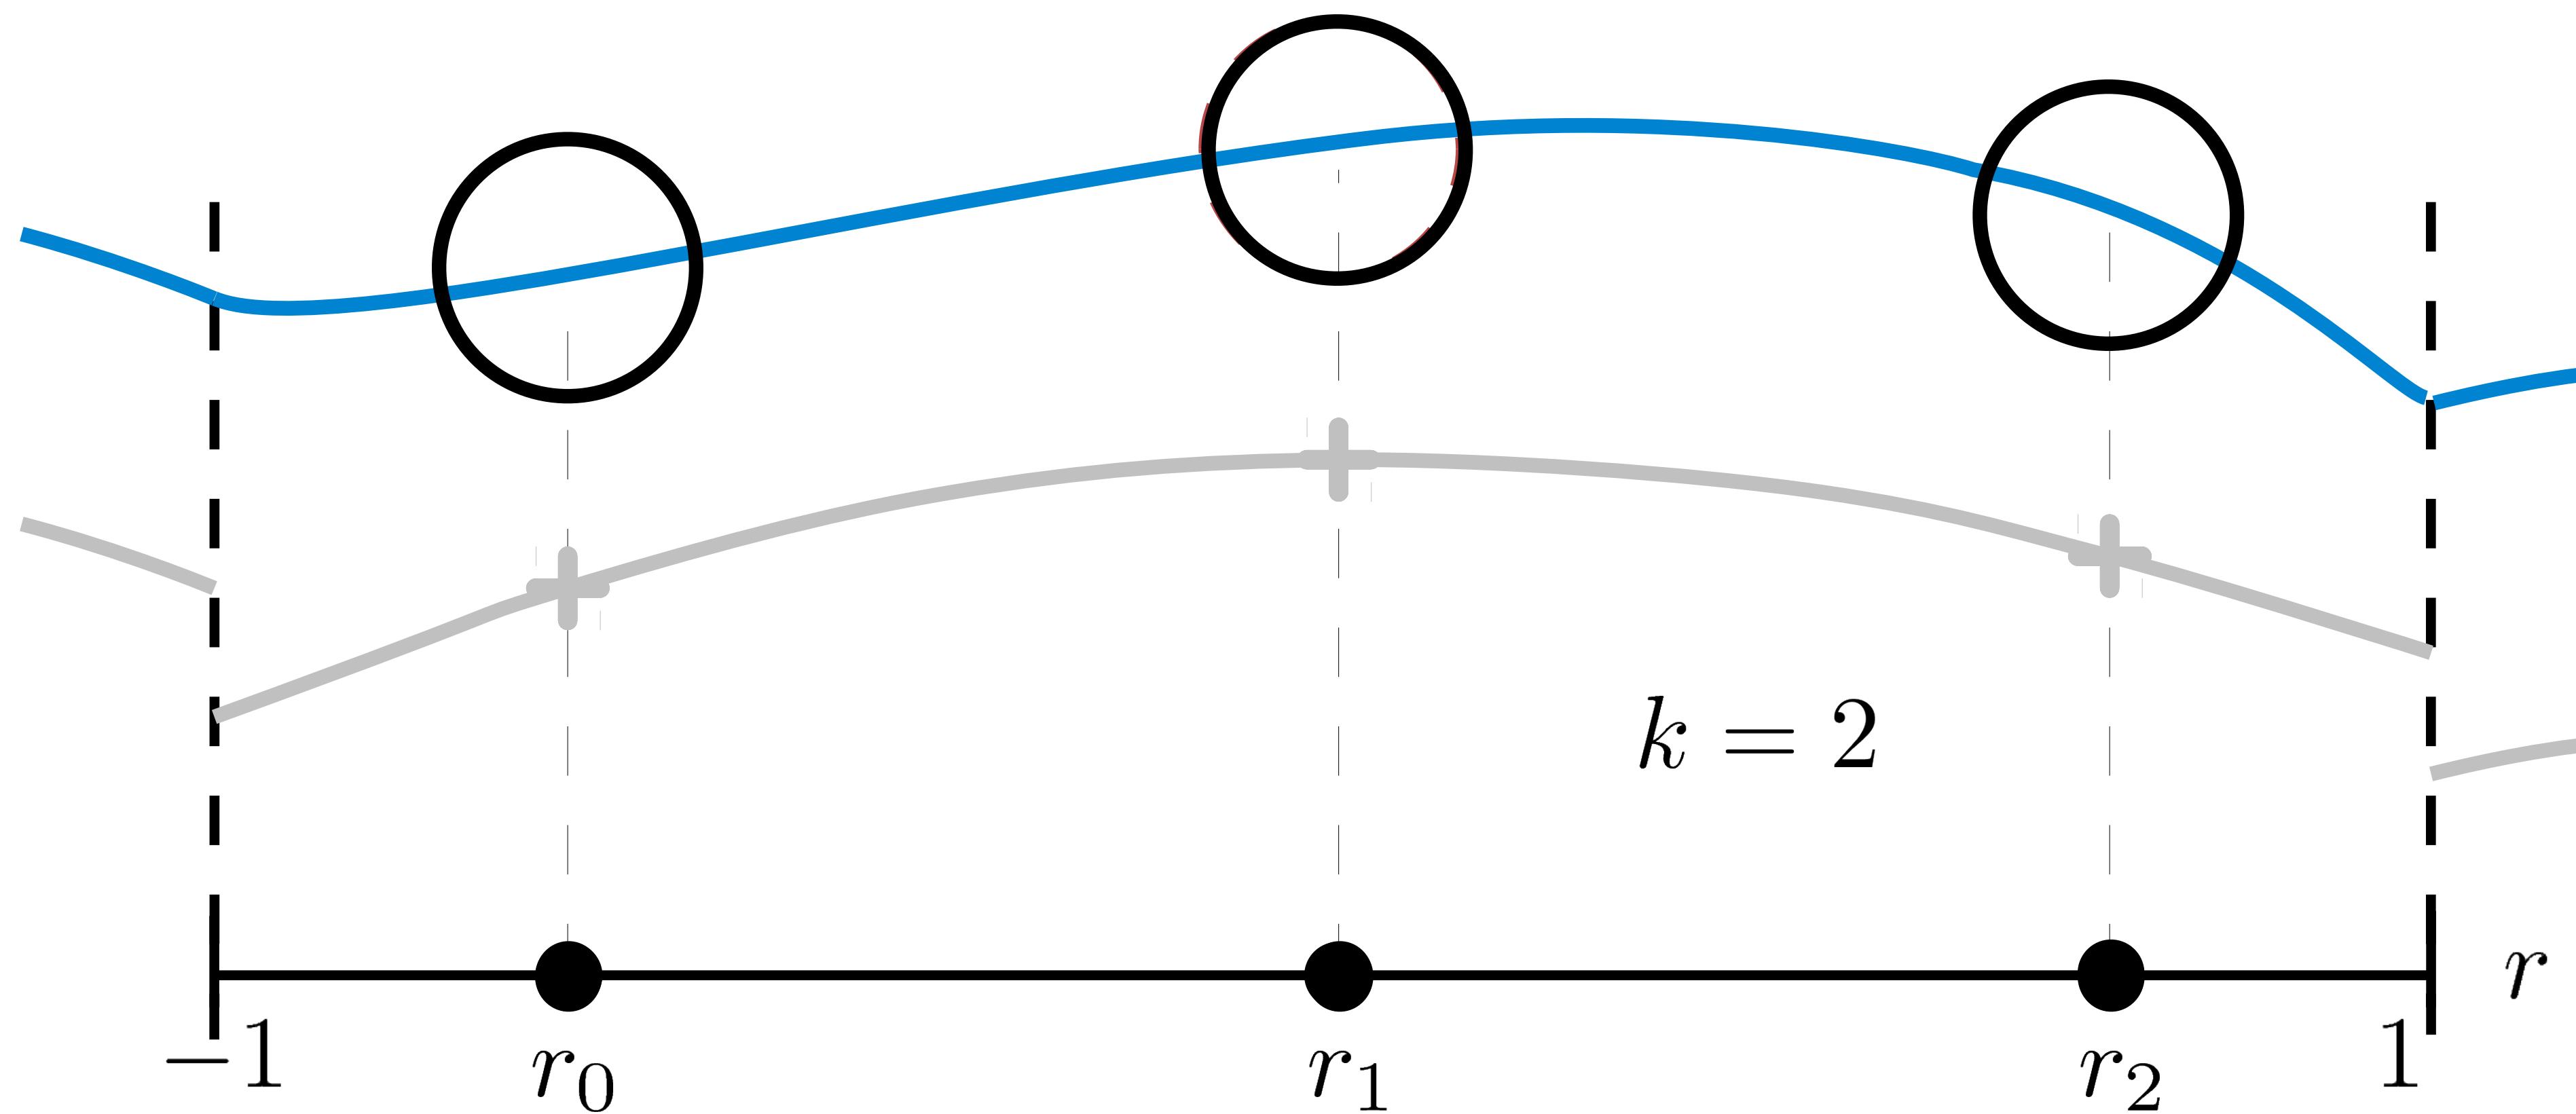
Theory

- Repeat from the right...



Theory

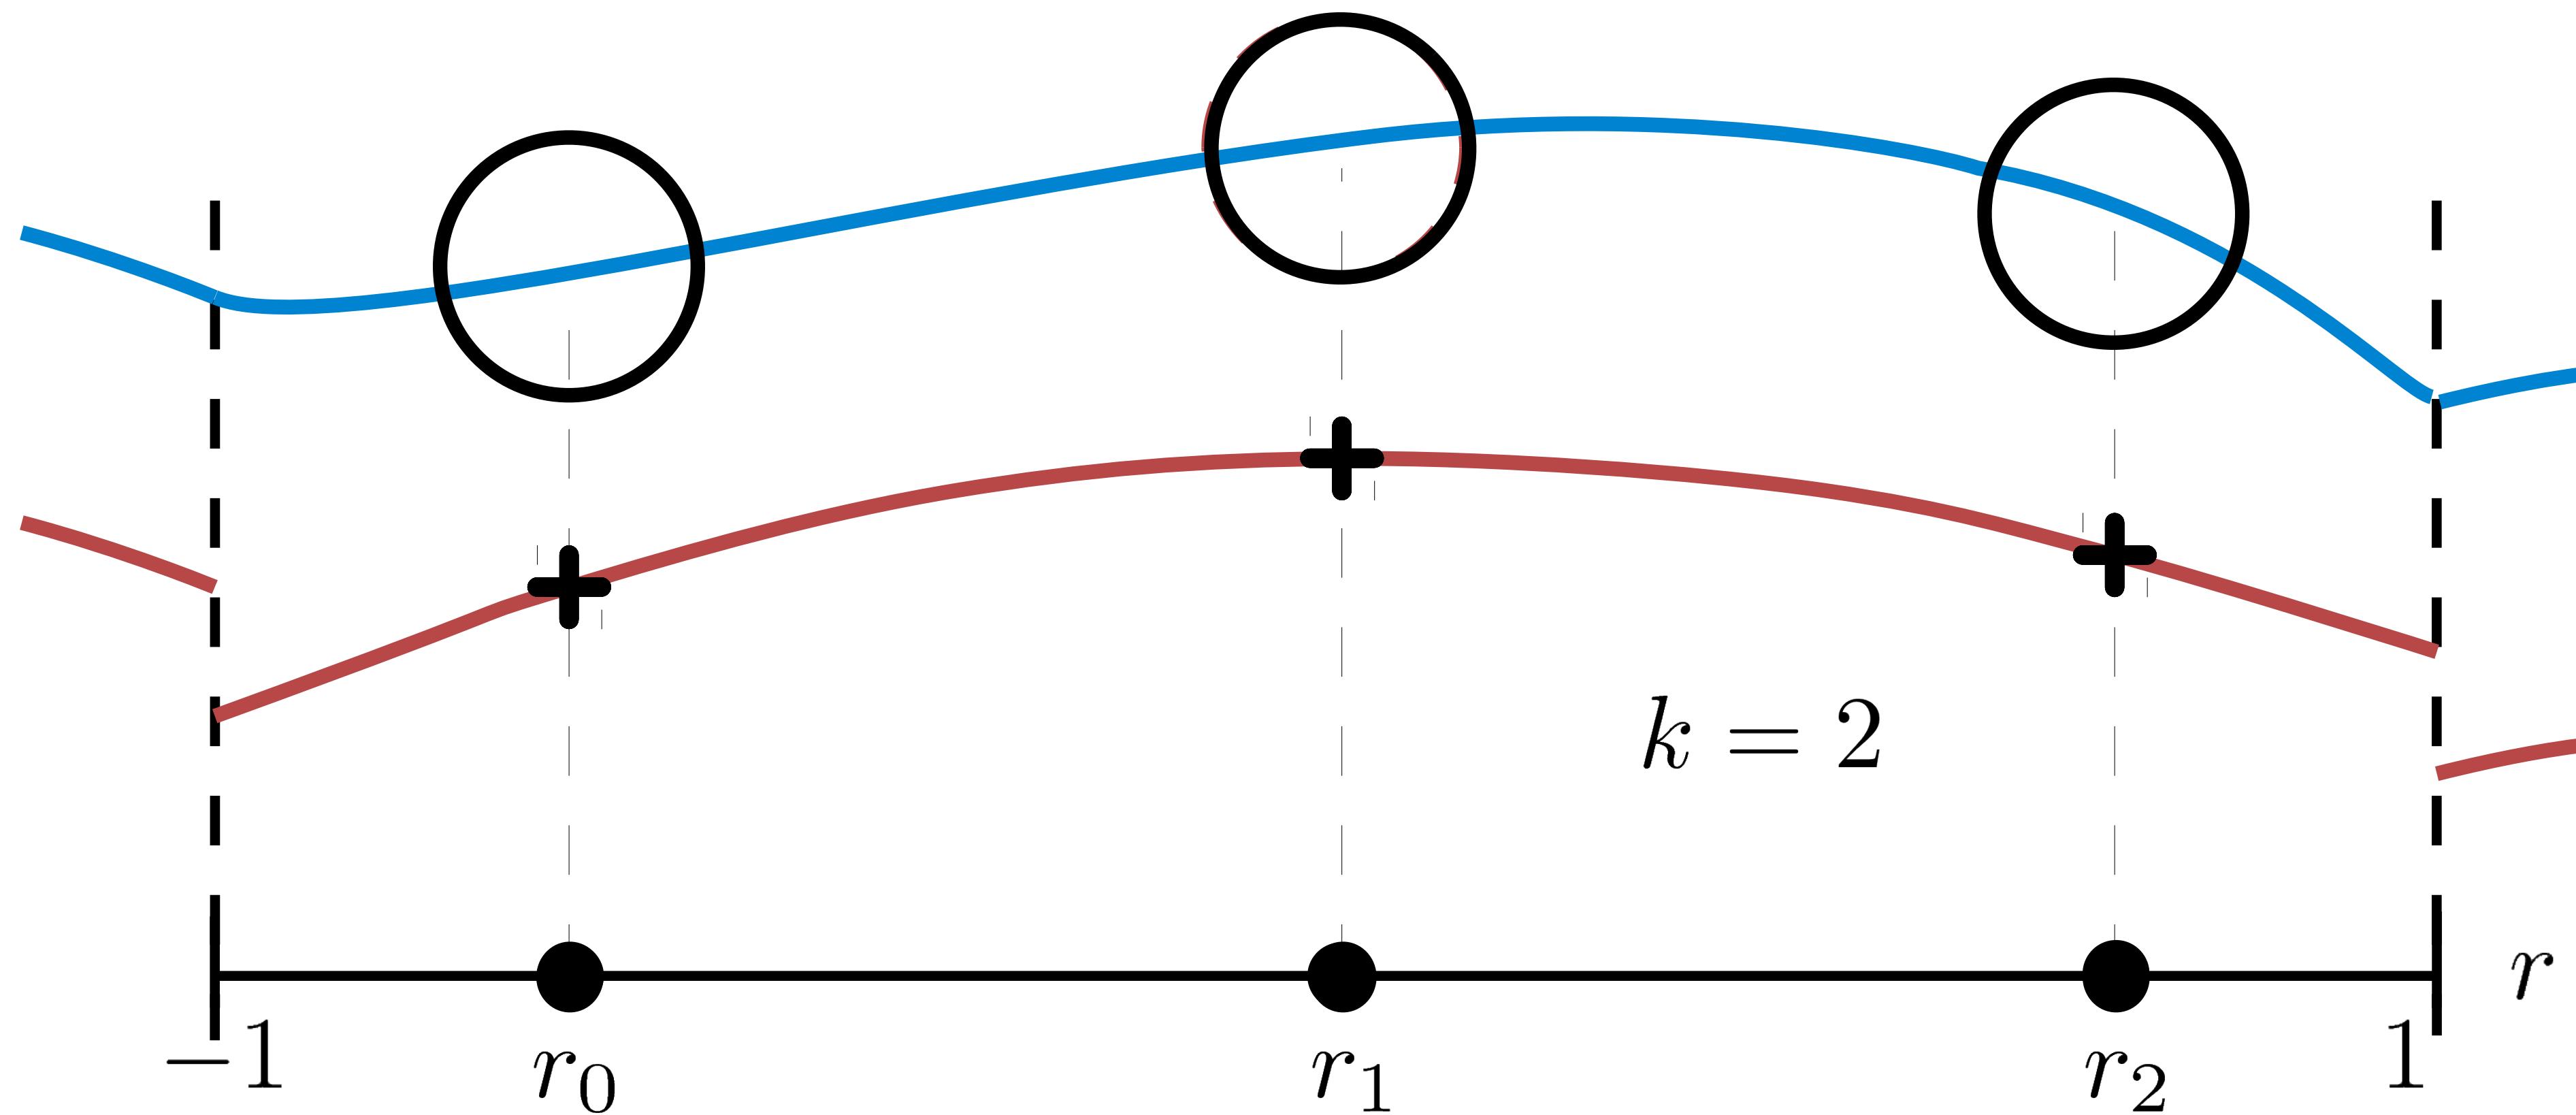
- Evaluate divergence of the continuous flux at the solution points



Theory

- And advance the solution in time...

$$\frac{\partial u}{\partial t} + \frac{\partial f}{\partial x} = 0 \quad \rightarrow \quad \frac{\partial \hat{u}_i^\delta}{\partial t} = - \frac{\partial \hat{f}^\delta}{\partial r}(r_i)$$



Theory

- Nature of FR scheme depends on location of solution points, interface flux, correction function
- Can recover nodal Discontinuous Galerkin (DG) method, and any Spectral Difference (SD) method via judicious choice of correction function [I]

[I] H.T. Huynh. A flux reconstruction approach to high-order schemes including discontinuous Galerkin methods. AIAA Paper 2007-4079. 2007

Theory

- In 2011 Vincent, Castonguay and Jameson identified a range of stable correction functions for all orders of accuracy using an energy method [2][3]

- [2] P. E. Vincent, P. Castonguay, A. Jameson. A New Class of High-Order Energy Stable Flux Reconstruction Schemes. *Journal of Scientific Computing*. 2011
- [3] A. Jameson, P. E. Vincent, P. Castonguay. On the Non-Linear Stability of High-Order Flux Reconstruction Schemes. *Journal of Scientific Computing*. 2011

Theory

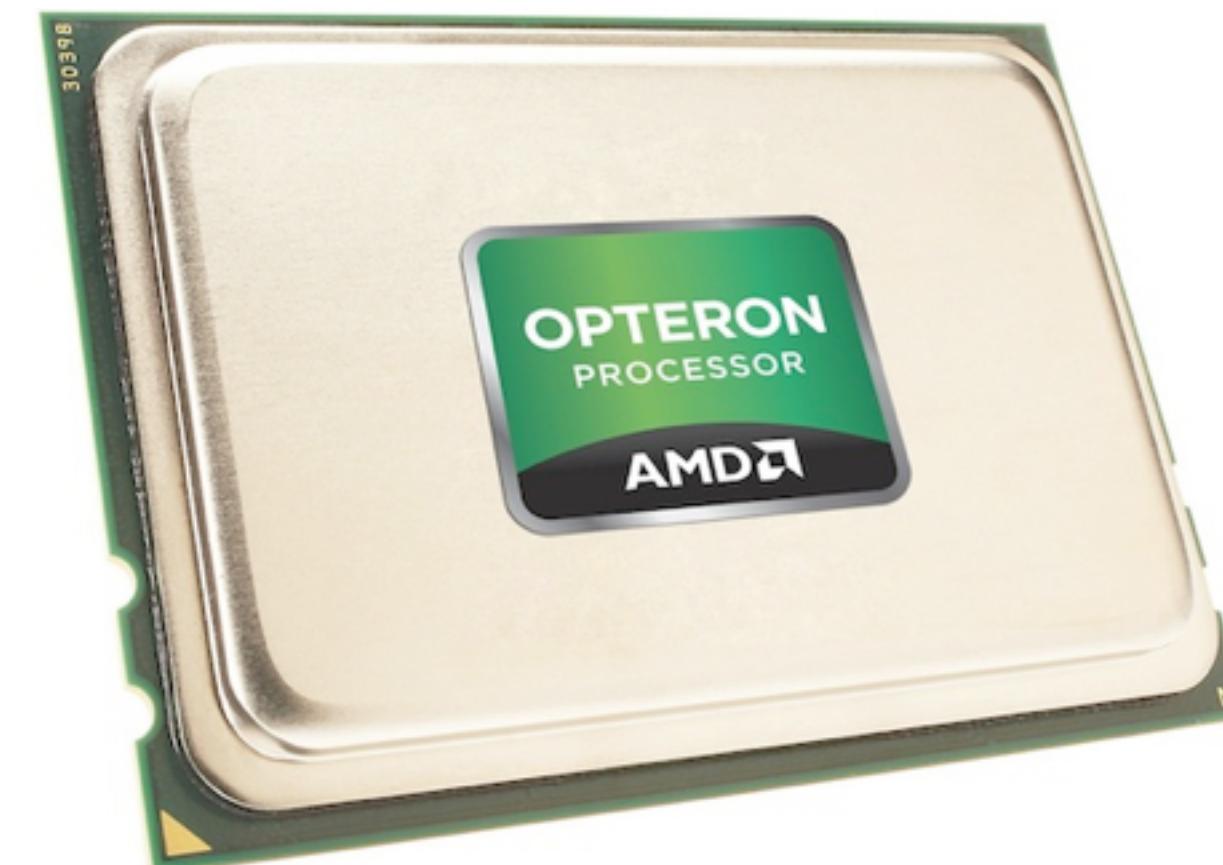
- Most operations point/element local
- No explicit quadrature (c.f. DG)
- No global assembly (c.f. FEM)
- No matrix inversions

PyFR

- Recent rapid evolution of hardware available for scientific computing – exciting times

PyFR

- Multi-core CPUs now ubiquitous (with 10's of cores per device)
- Rely extensively on vectorization for throughput



PyFR

- Many-core accelerators a hot topic (100s - 1000s cores per device)

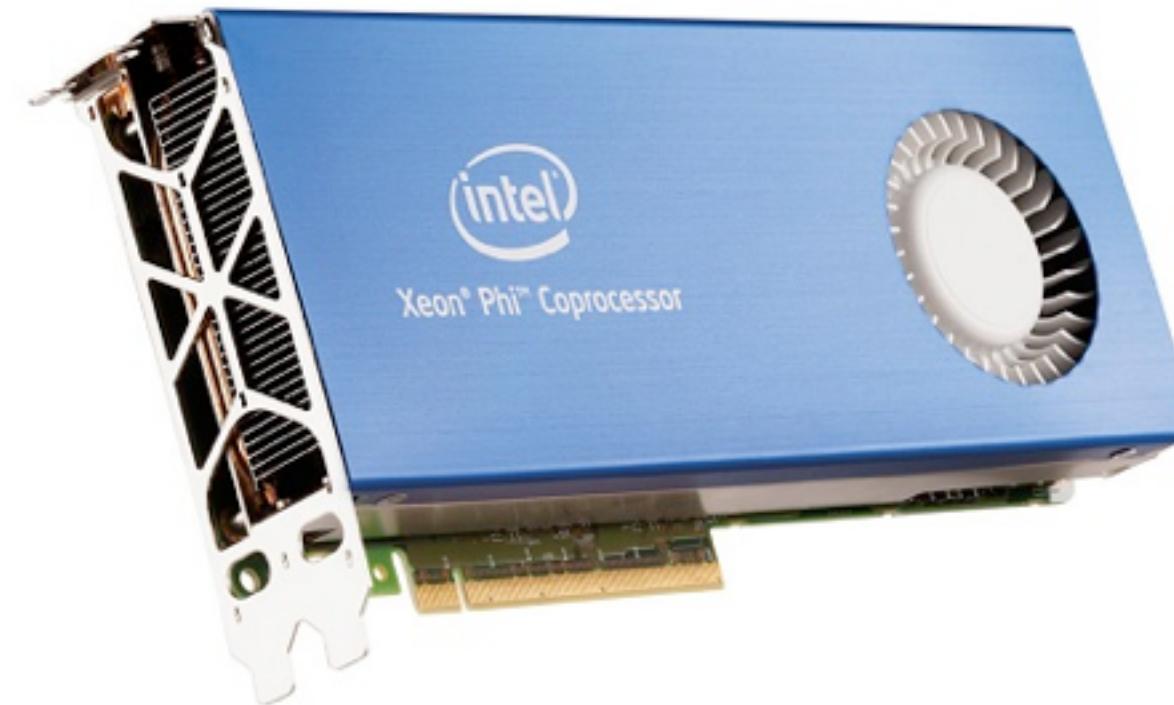


PyFR

- GPUs may or may not be the future... but many-core + low memory-per-core quite likely is



Nvidia K20



Intel Phi



AMD SI10000

PyFR

- What approach have we adopted?
- Write platform-independent code in Python and interface with accelerator via bespoke kernels
- Leverage the near universal availability of optimized BLAS libraries for CPUs/accelerators

PyFR

- Current status

Capabilities

- Compressible 2D/3D Navier-Stokes
- Unstructured grids of quads or hexahedra
- High-order FR in space
- Explicit Runge-Kutta in time

Implementation

- Setup (SymPy + NumPy)
- Kernel generation (Mako)
- Kernel invocation (PyCUDA &c.)
- Communication (mpi4py)

```
graph TD; Cap[Capabilities] --> CUDA[CUDA]; Impl[Implementation] --> COpenMP[C/OpenMP]
```

CUDA

C/OpenMP

PyFR

- Current status
- 4000 lines of Python
- 800 lines of Mako/CUDA
- 800 lines of Mako/C
- Slated for release under a BSD license in late summer

PyFR

- Theory of VCJH correction functions

$$\eta_k \in \left\{ 0, \frac{k}{k+1}, \frac{k+1}{k}, \dots \right\} \quad g'_R = \frac{1}{2} \frac{d}{dr} \left[L_k + \frac{\eta_k L_{k-1} + L_{k+1}}{1 + \eta_k} \right]$$

- Implementation using SymPy

```
def diff_vcjh_correctionfn(k, eta, sym):
    # Expand shorthand forms of eta_k for common schemes
    etacommon = dict(dg='0', sd='k/(k+1)', hu='(k+1)/k')

    eta_k = sy.S(etacommon.get(eta, eta), locals=dict(k=k))

    lkm1, lk, lkpl = [sy.legendre_poly(m, sym) for m in [k-1, k, k+1]]

    # Correction function derivatives, Eq. 3.46 and 3.47
    diffgr = (sy.S(1)/2 * (lk + (eta_k*lkm1 + lkpl)/(1 + eta_k))).diff()
    diffgl = -diffgr.subs(sym, -sym)

return diffgl, diffgr
```

PyFR

- Kernel generation using Mako templates

```
<%include file='idx_of.cu.mak' />

__global__ void
negdivconf(int nupts, int neles,
           ${dtype}* __restrict__ tdivtconf,
           const ${dtype}* __restrict__ rcpdjac,
           int ldt, int ldr)
{
    int eidx = blockIdx.x * blockDim.x + threadIdx.x;

    if (eidx <= neles)
    {
        for (int uidx = 0; uidx <= nupts; ++uidx)
        {
            ${dtype} s = -rcpdjac[IDX_OF(uidx, eidx, ldr)];

            % for i in range(nvars):
                tdivtconf[U_IDX_OF(uidx, eidx, ${i}, neles, ldt)] *= s;
            % endfor
        }
    }
}
```

PyFR

- Kernel generation using Mako templates

```
#ifndef _PYFR_IDX_OF
#define _PYFR_IDX_OF

#define IDX_OF(i, j, ldim) ((i)*(ldim) + (j))

#define U_IDX_OF(upt, ele, var, nele, ldim) \
    IDX_OF(upt, nele*var + ele, ldim)

#endif // _PYFR_IDX_OF

__global__ void
negdivconf(int nupts, int neles,
           double* __restrict__ tdivtconf,
           const double* __restrict__ rcpdjac,
           int ldt, int ldr)
{
    int eidx = blockIdx.x * blockDim.x + threadIdx.x;

    if (eidx < neles)
    {
        for (int uidx = 0; uidx < nupts; ++uidx)
        {
            double s = -rcpdjac[IDX_OF(uidx, eidx, ldr)];

            tdivtconf[U_IDX_OF(uidx, eidx, 0, neles, ldt)] *= s;
            tdivtconf[U_IDX_OF(uidx, eidx, 1, neles, ldt)] *= s;
            tdivtconf[U_IDX_OF(uidx, eidx, 2, neles, ldt)] *= s;
            tdivtconf[U_IDX_OF(uidx, eidx, 3, neles, ldt)] *= s;
            tdivtconf[U_IDX_OF(uidx, eidx, 4, neles, ldt)] *= s;
        }
    }
}
```

PyFR

- Queued compute and MPI kernel execution

```
def _get_negdivf(self):  
    runall = self._backend.runall  
    q1, q2 = self._queues  
  
    q1 << self._disu_fpts_kerns()  
    q1 << self._mpi_inters_scal_fpts0_pack_kerns()  
    runall([q1])  
  
    q1 << self._tdisf_upts_kerns()  
    q1 << self._tdivtpcorf_upts_kerns()  
    q1 << self._int_inters_rsolve_kerns()  
    q1 << self._bc_inters_rsolve_kerns()  
  
    q2 << self._mpi_inters_scal_fpts0_send_kerns()  
    q2 << self._mpi_inters_scal_fpts0_recv_kerns()  
    q2 << self._mpi_inters_scal_fpts0_unpack_kerns()  
  
    runall([q1, q2])  
  
    q1 << self._mpi_inters_rsolve_kerns()  
    q1 << self._tdivtconf_upts_kerns()  
    q1 << self._negdivconf_upts_kerns()  
    runall([q1])
```

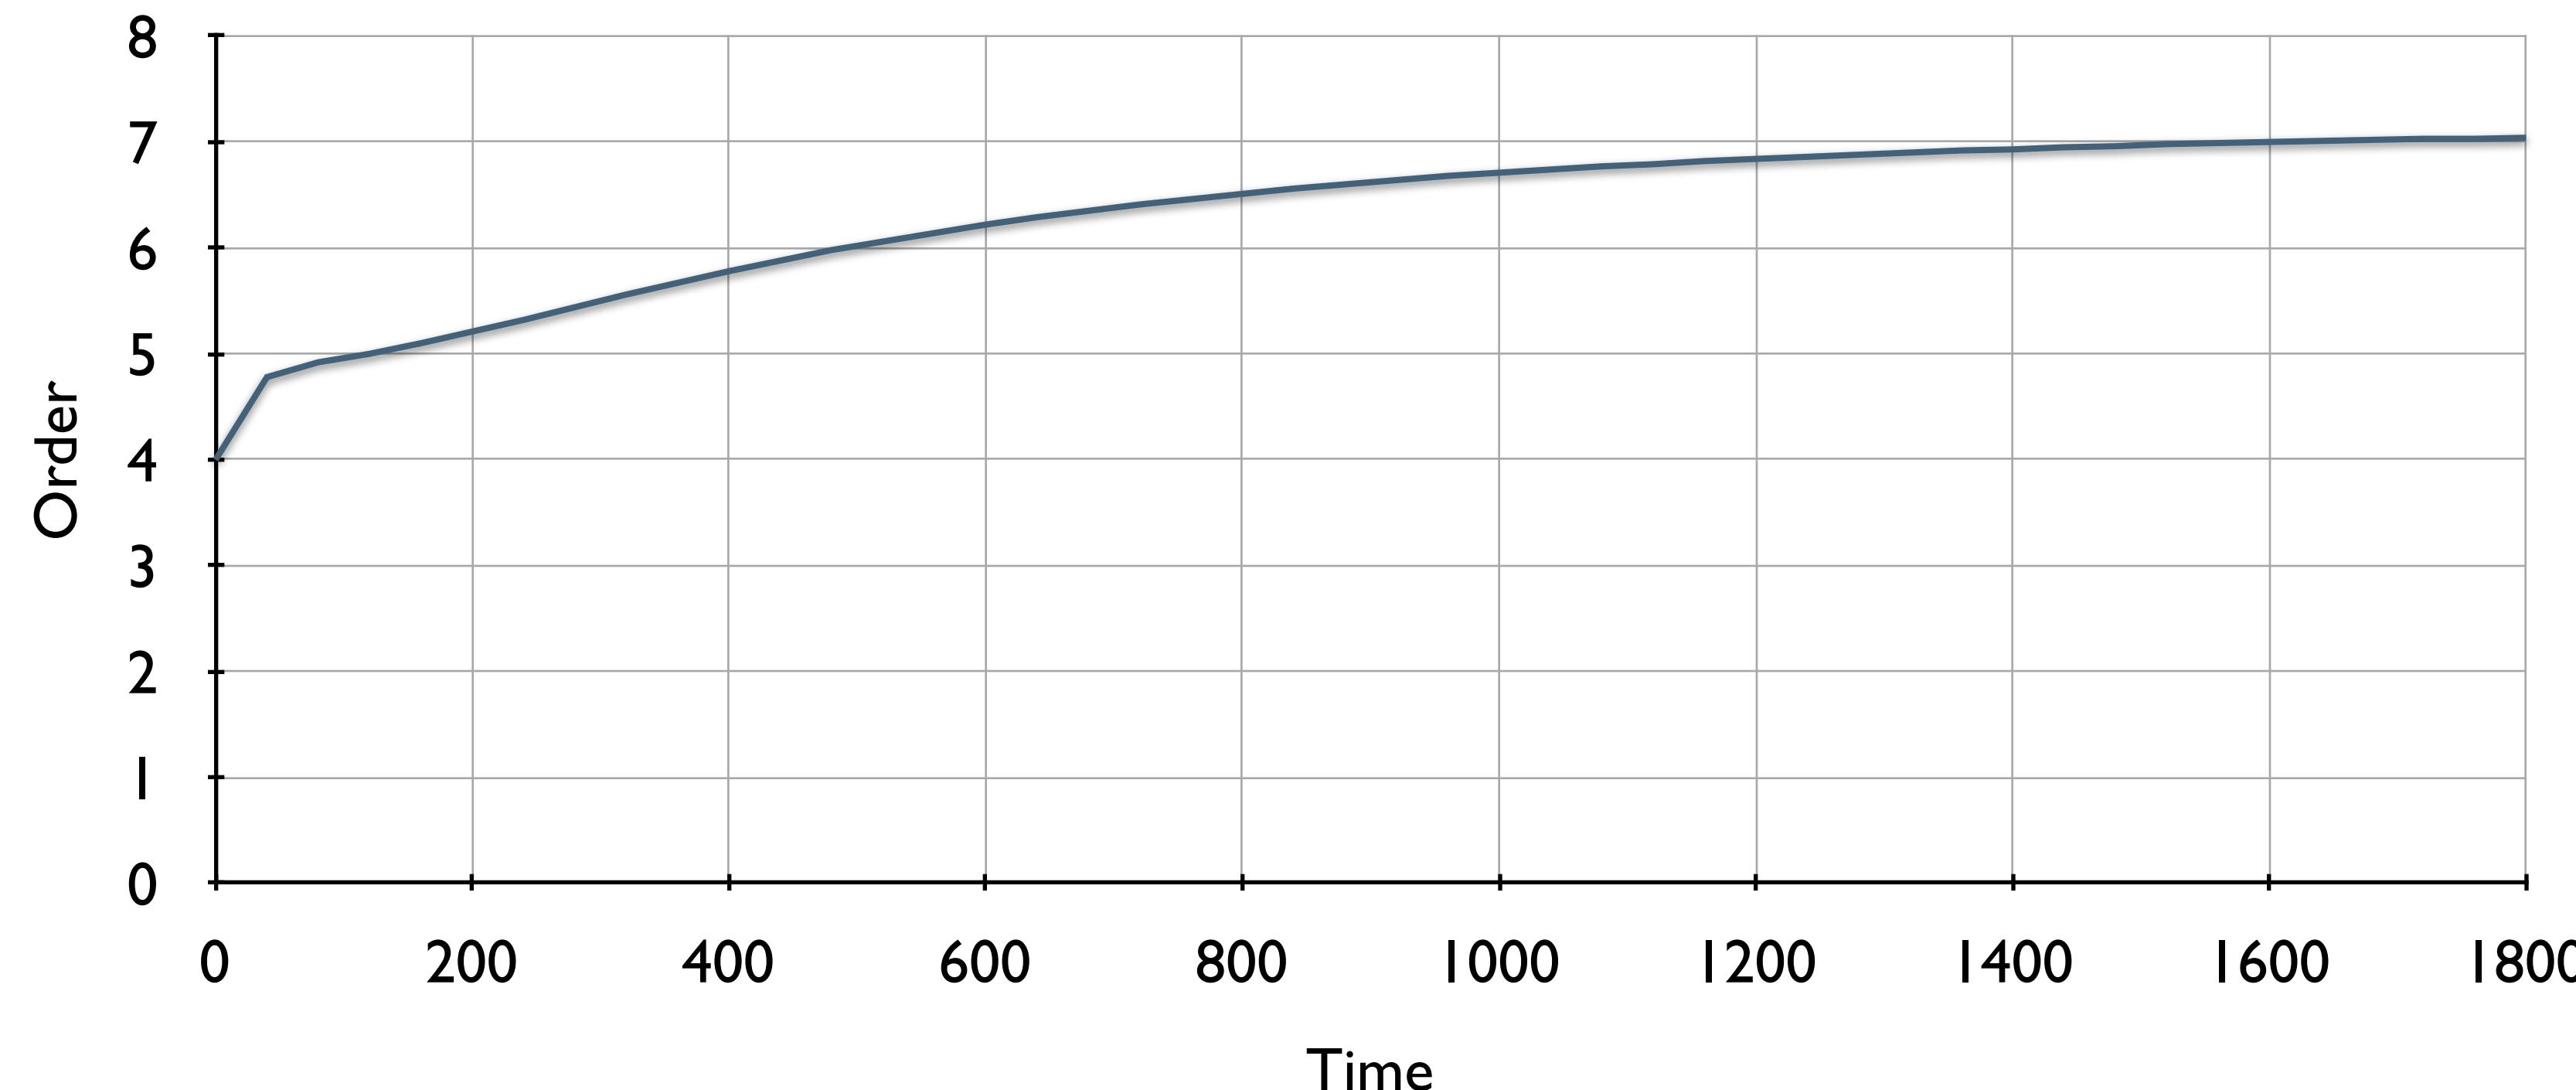
Results

- Some preliminary results on our $4 \times \text{K20}$ workstation



Results

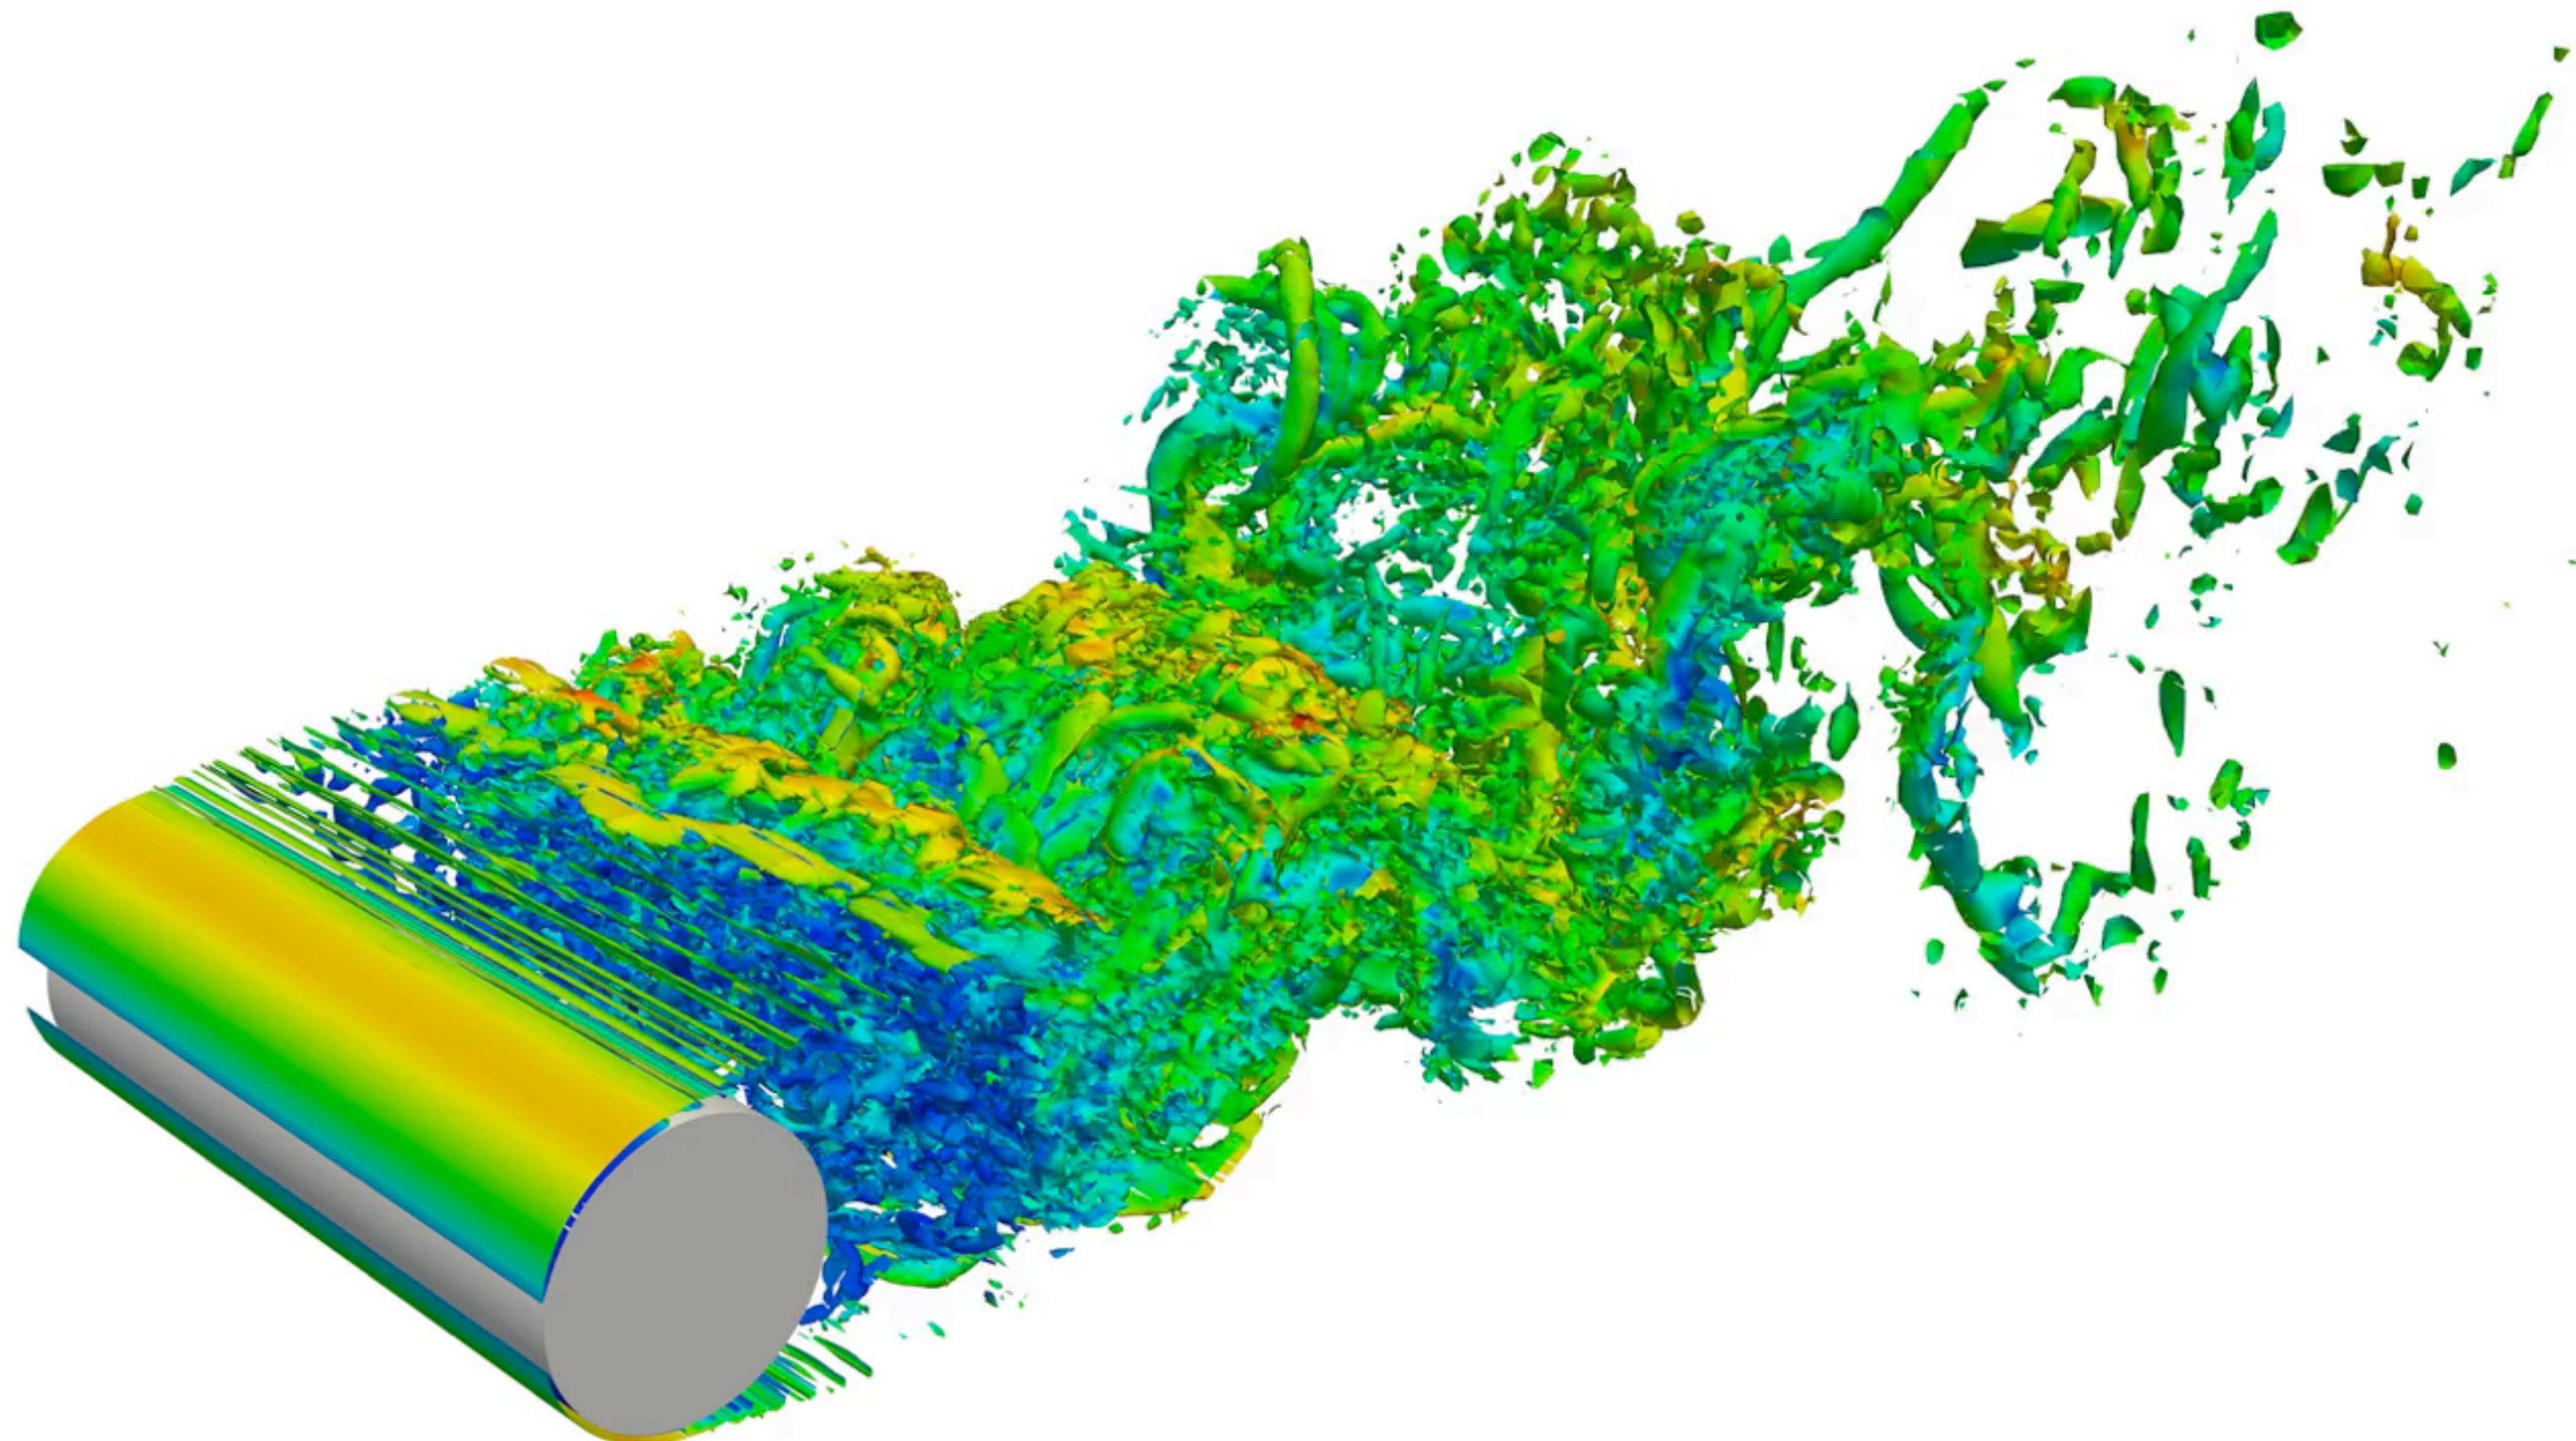
- Super $2k + 1$ accuracy for a 2D Euler vortex with $k = 3$ [4]



[4] P.E.Vincent, P.Castonguay, A.Jameson. Insights from von Neumann analysis of high-order flux reconstruction schemes. Journal of Computational Physics. 2011

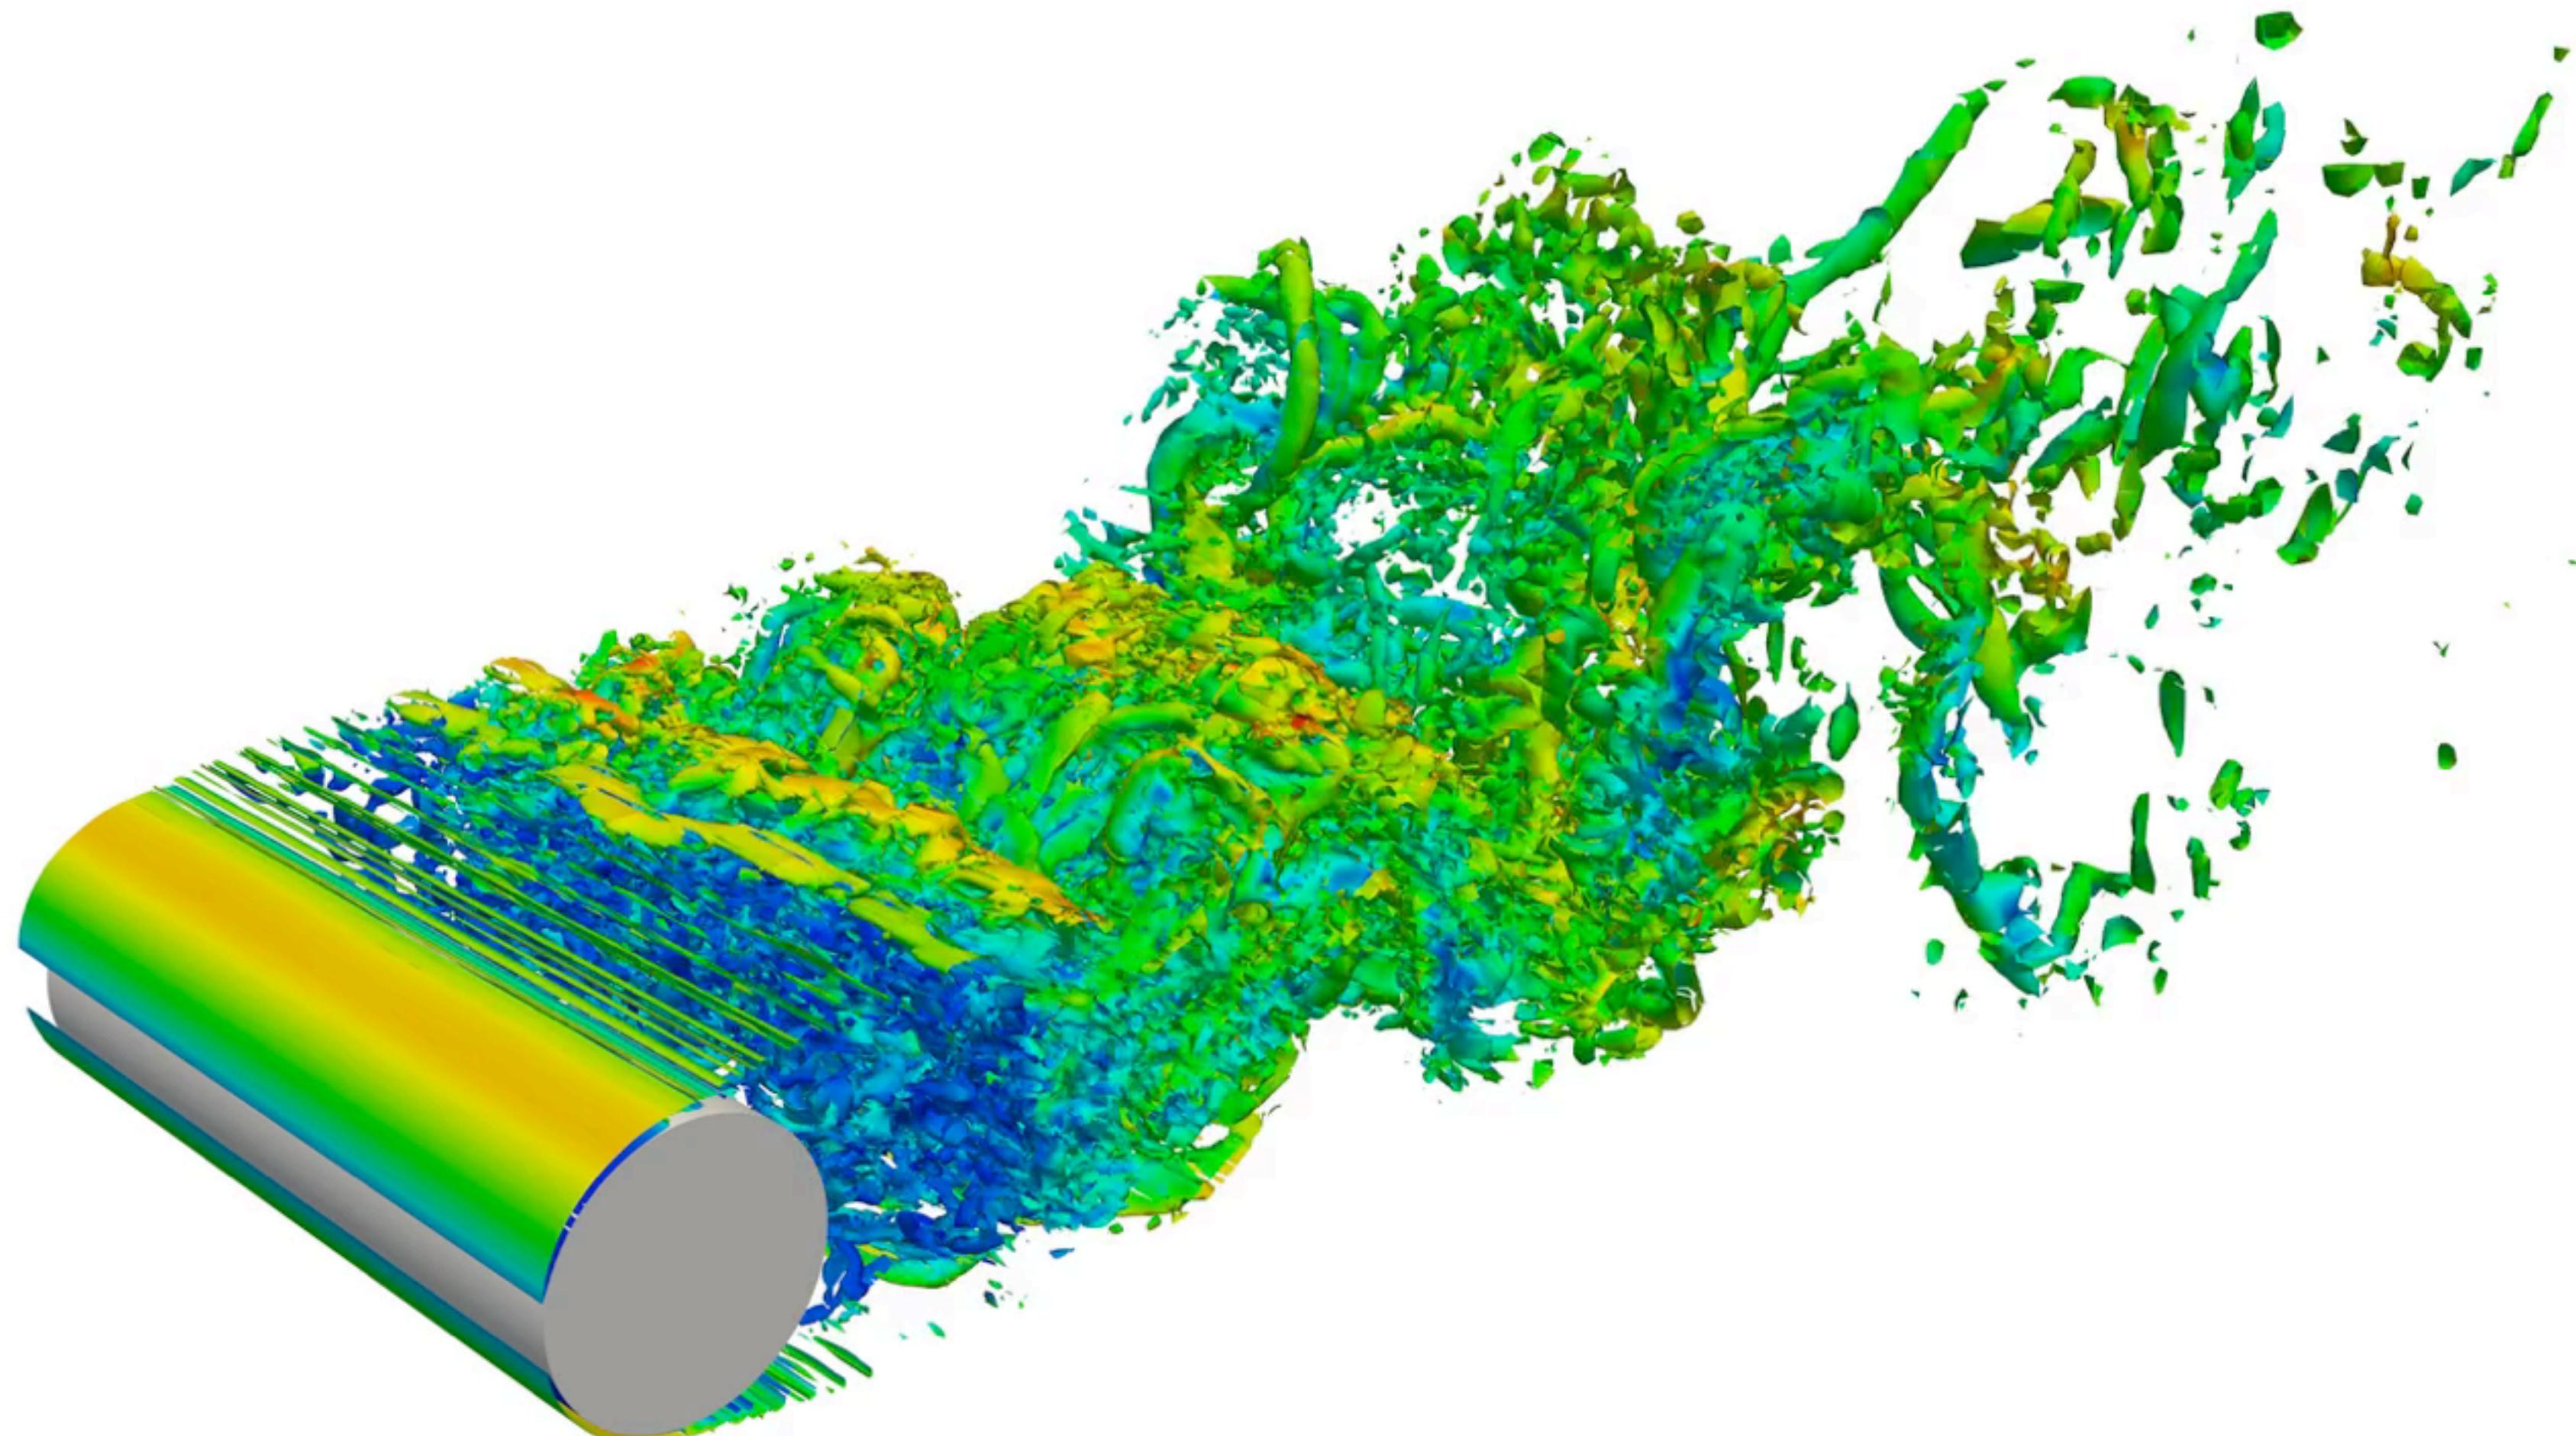
Results

- Flow over cylinder at $\text{Ma} = 0.2$ $\text{Re} = 3900$



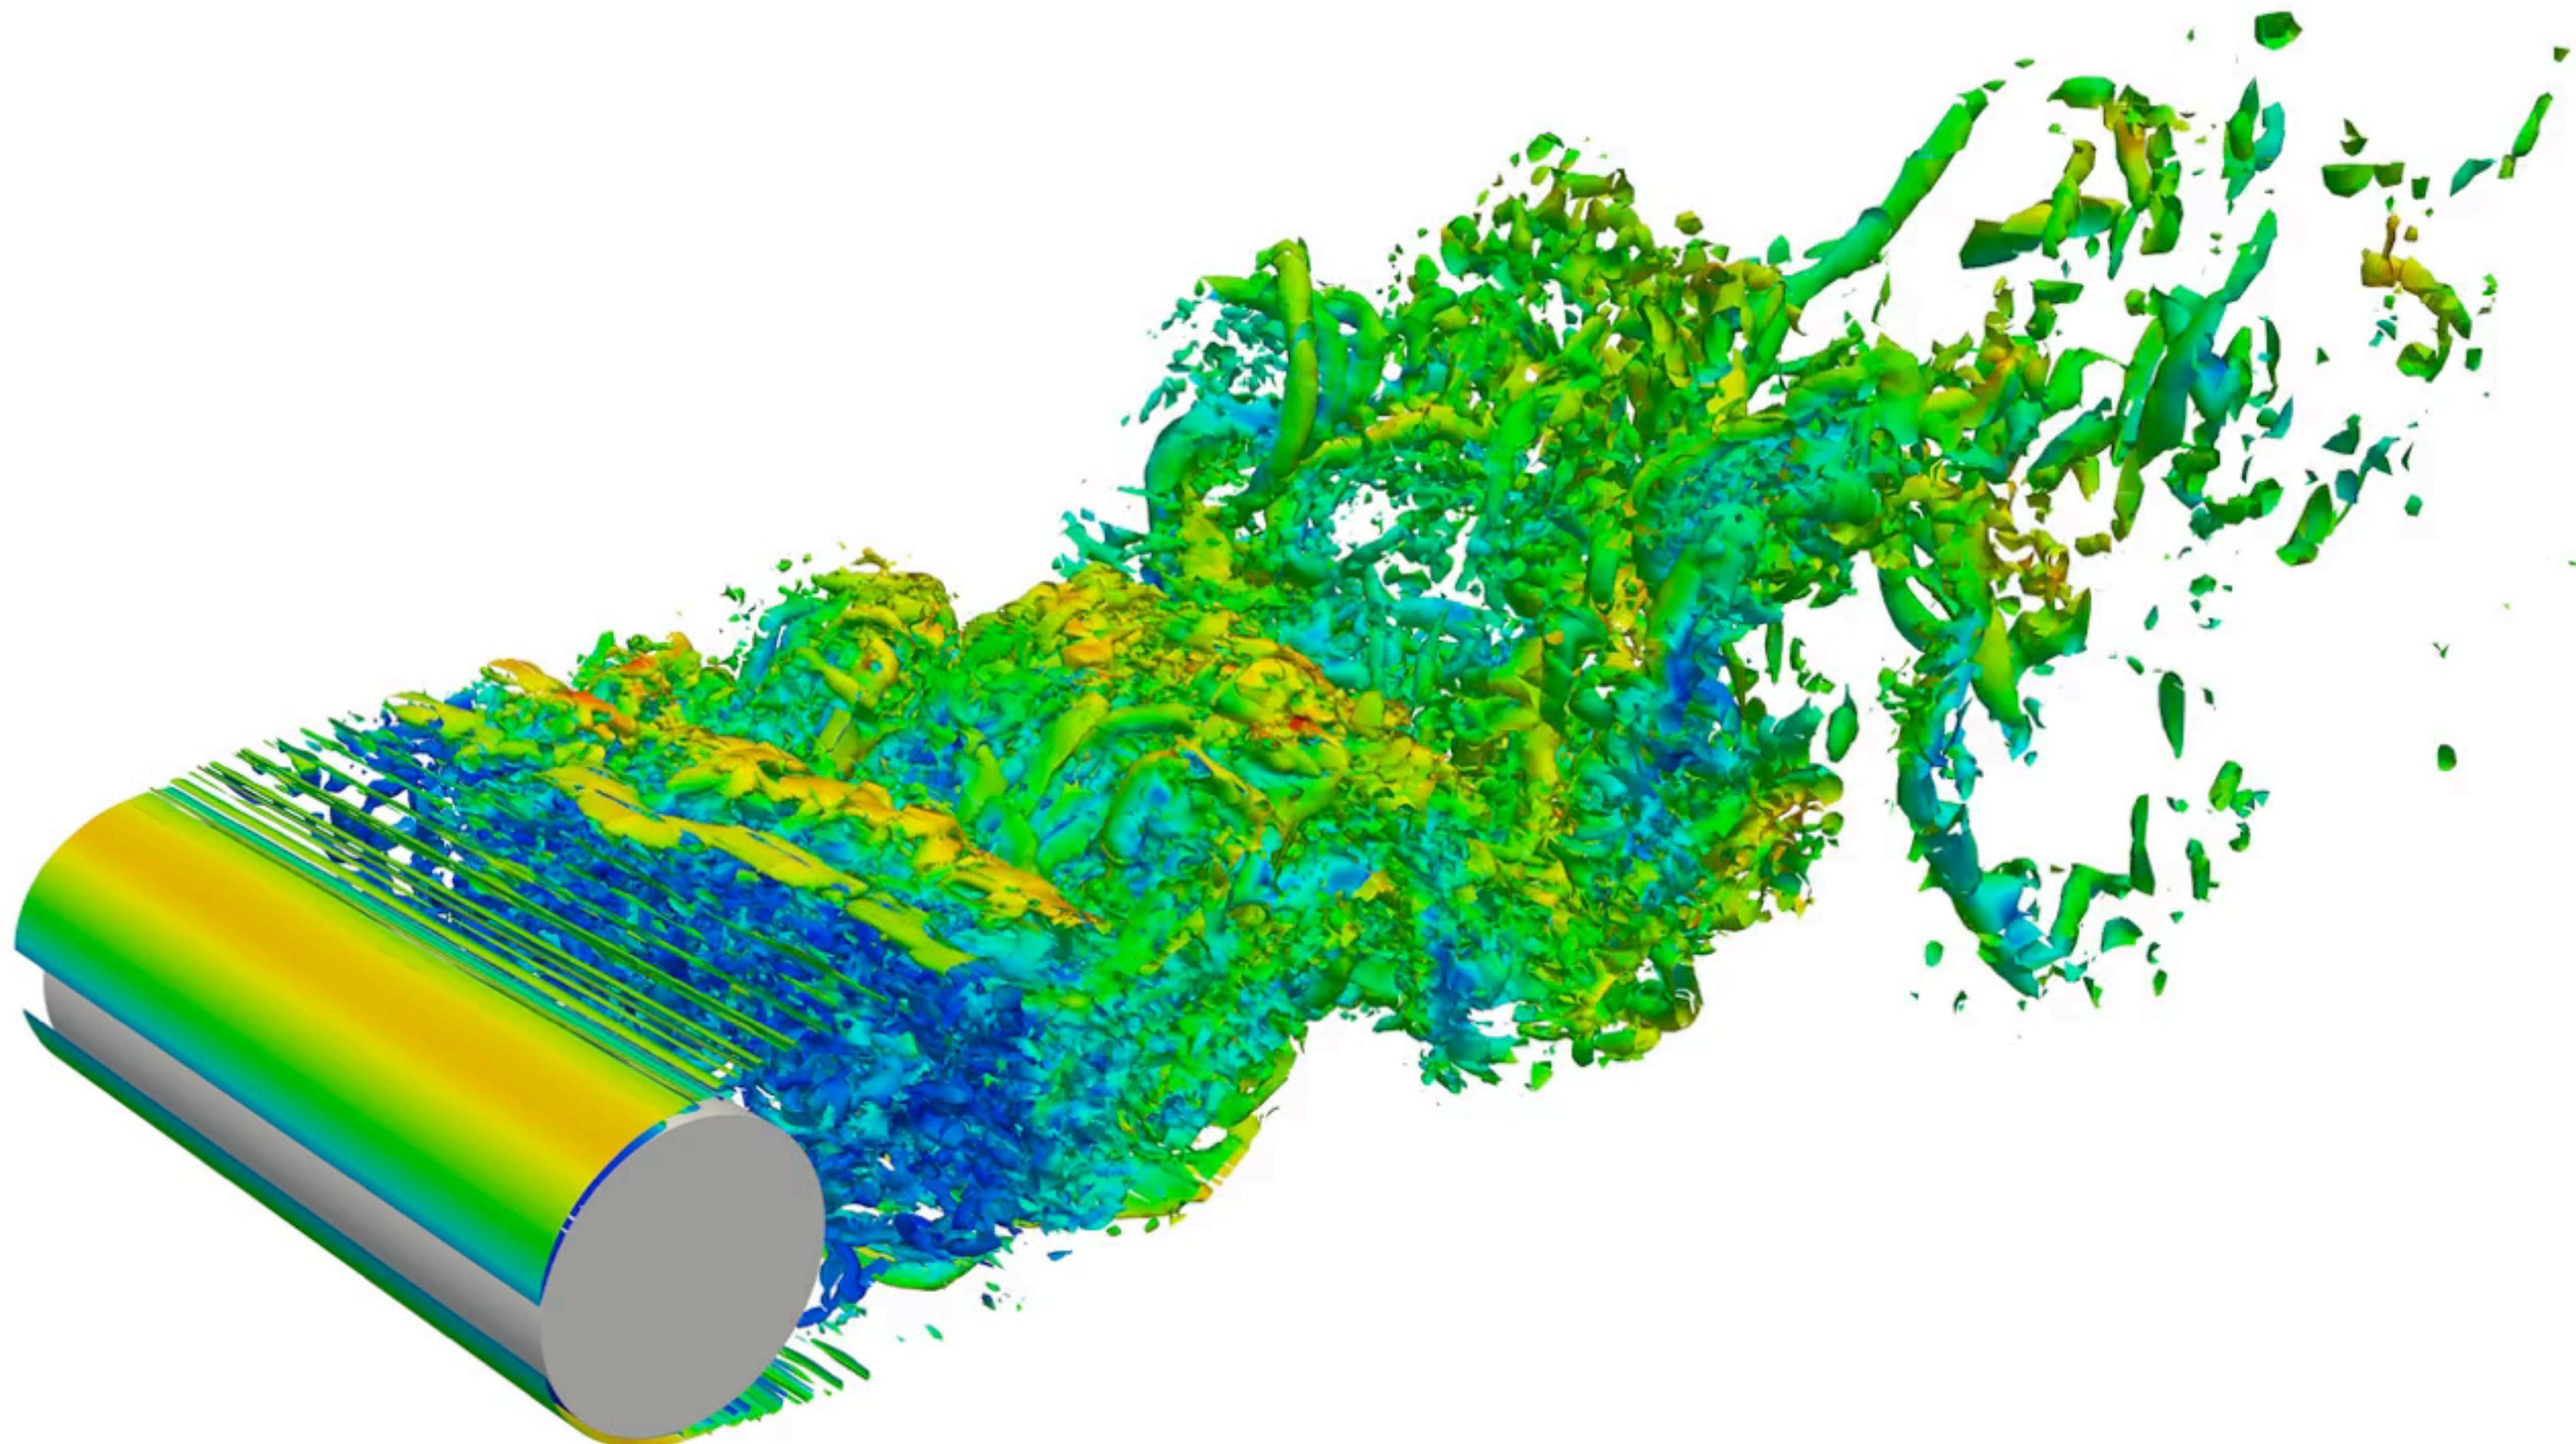
Results

- 5th order accurate with 2.9×10^7 DOFs



Results

- Double precision



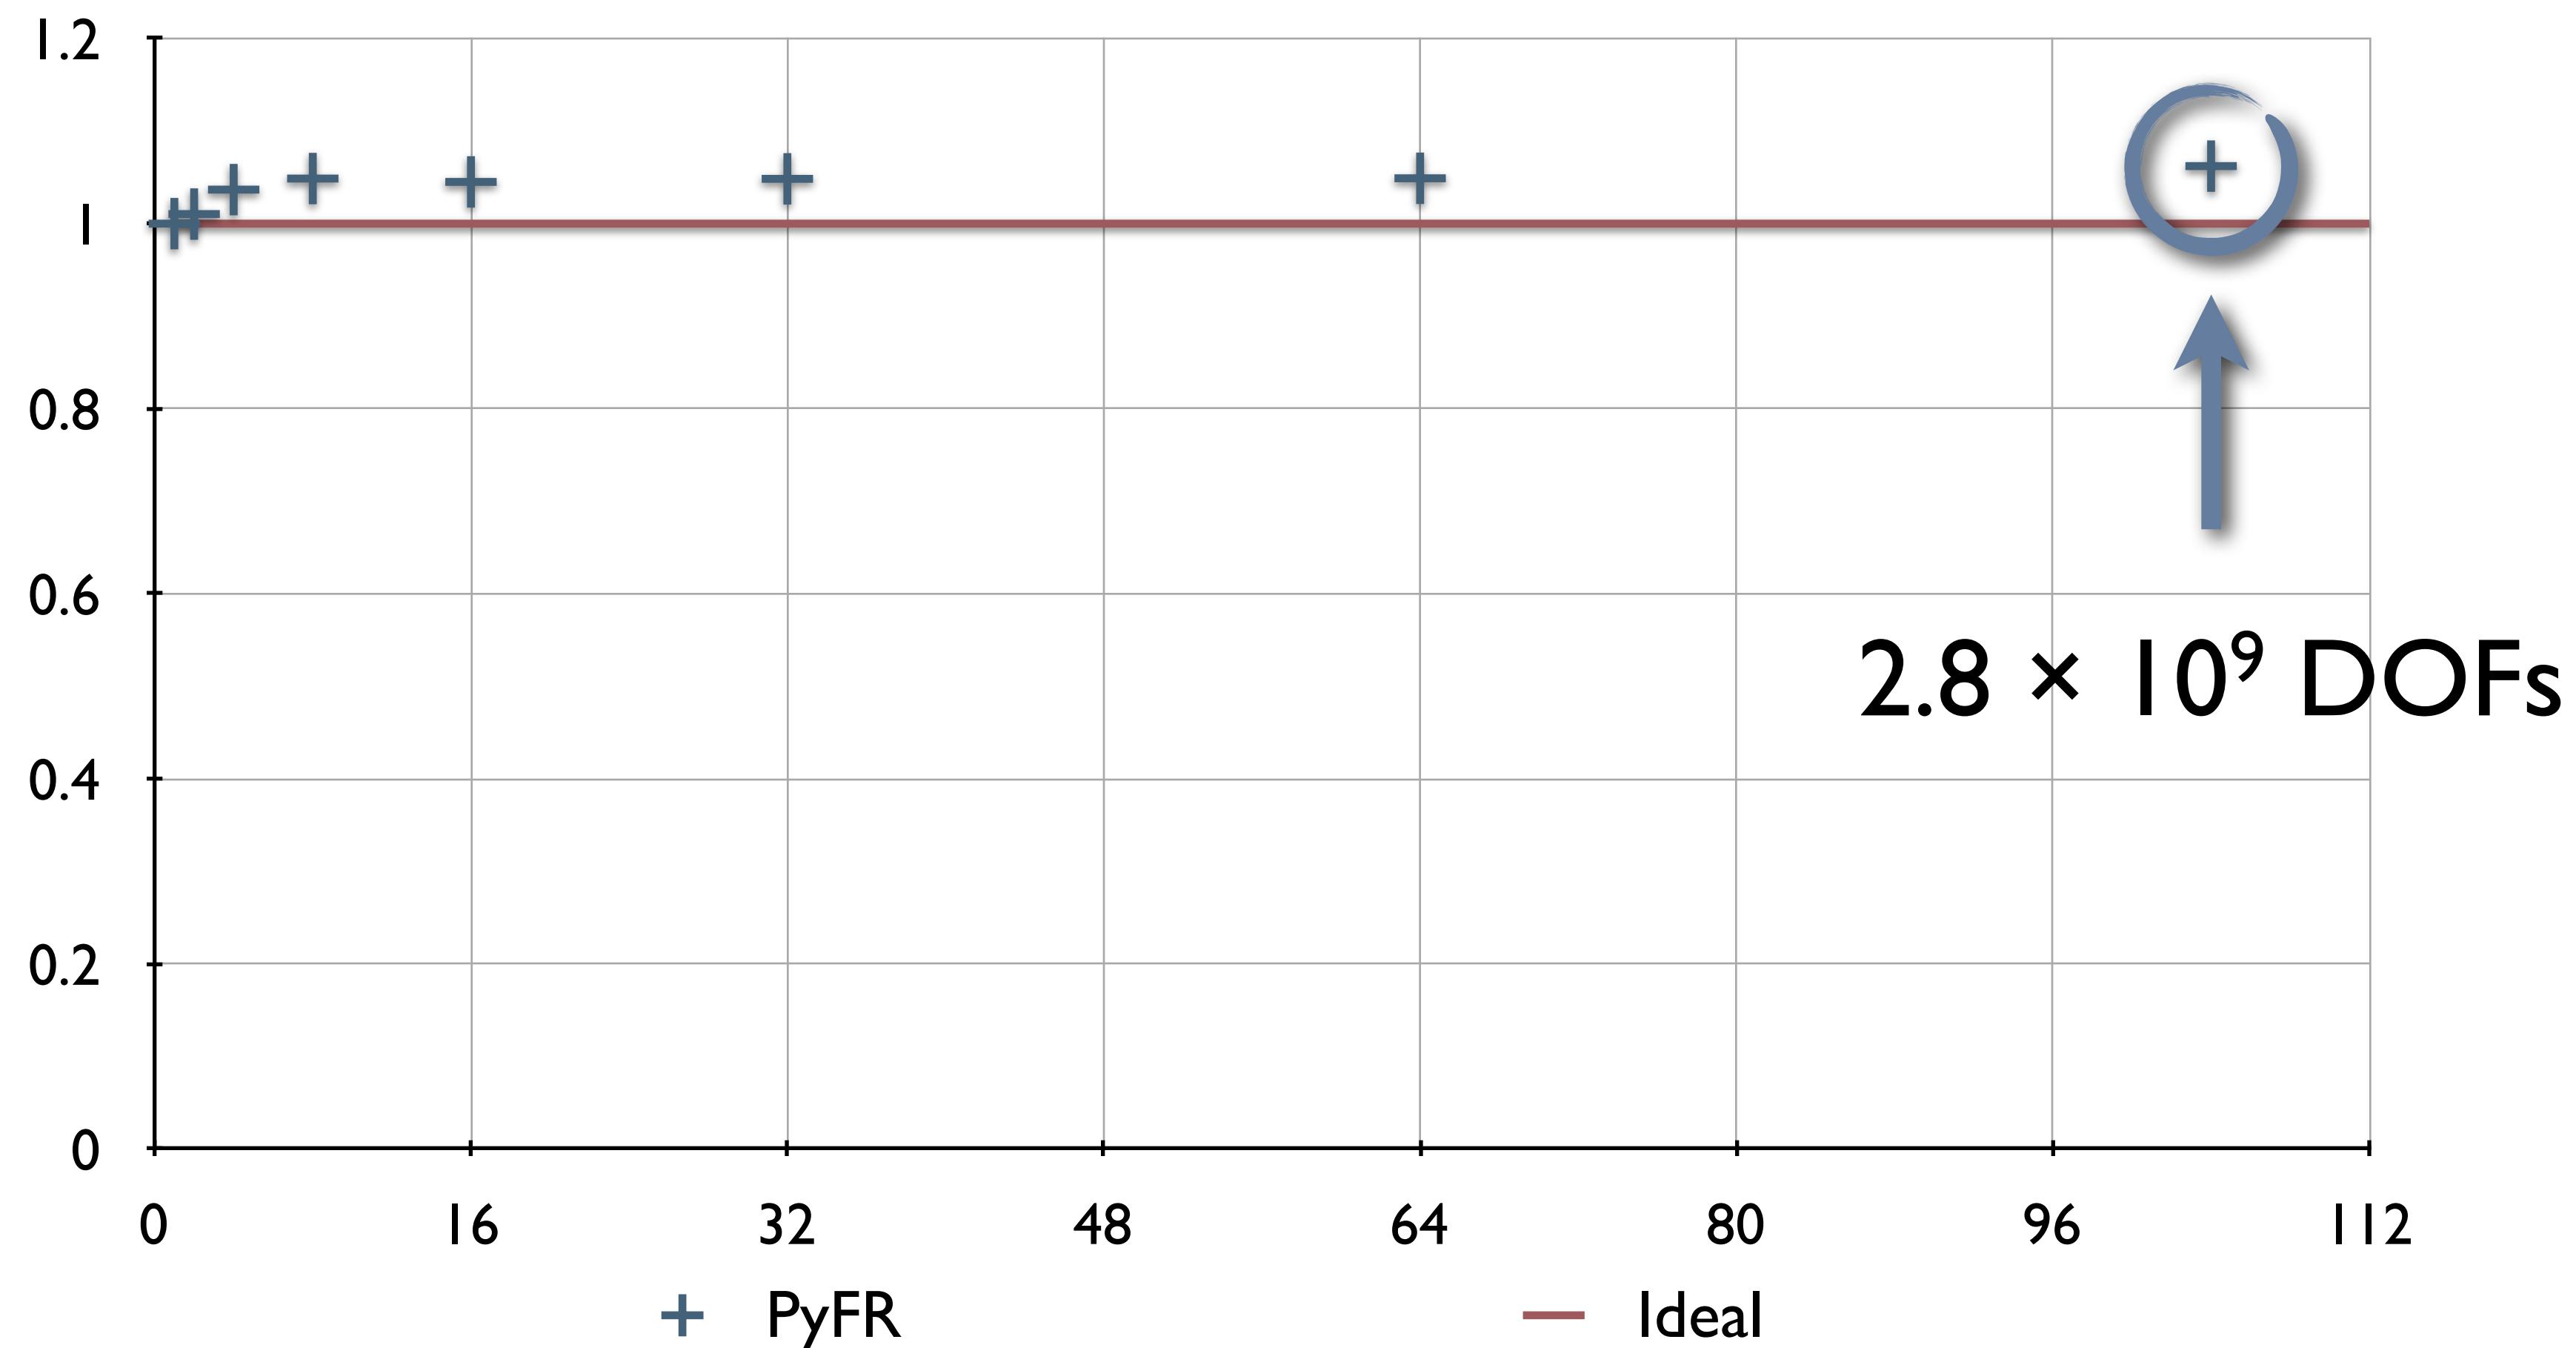
Results

- Scaling on **Emerald**, UK's largest multi-GPU cluster
- $372 \times$ Nvidia M2090s each with 512 CUDA cores



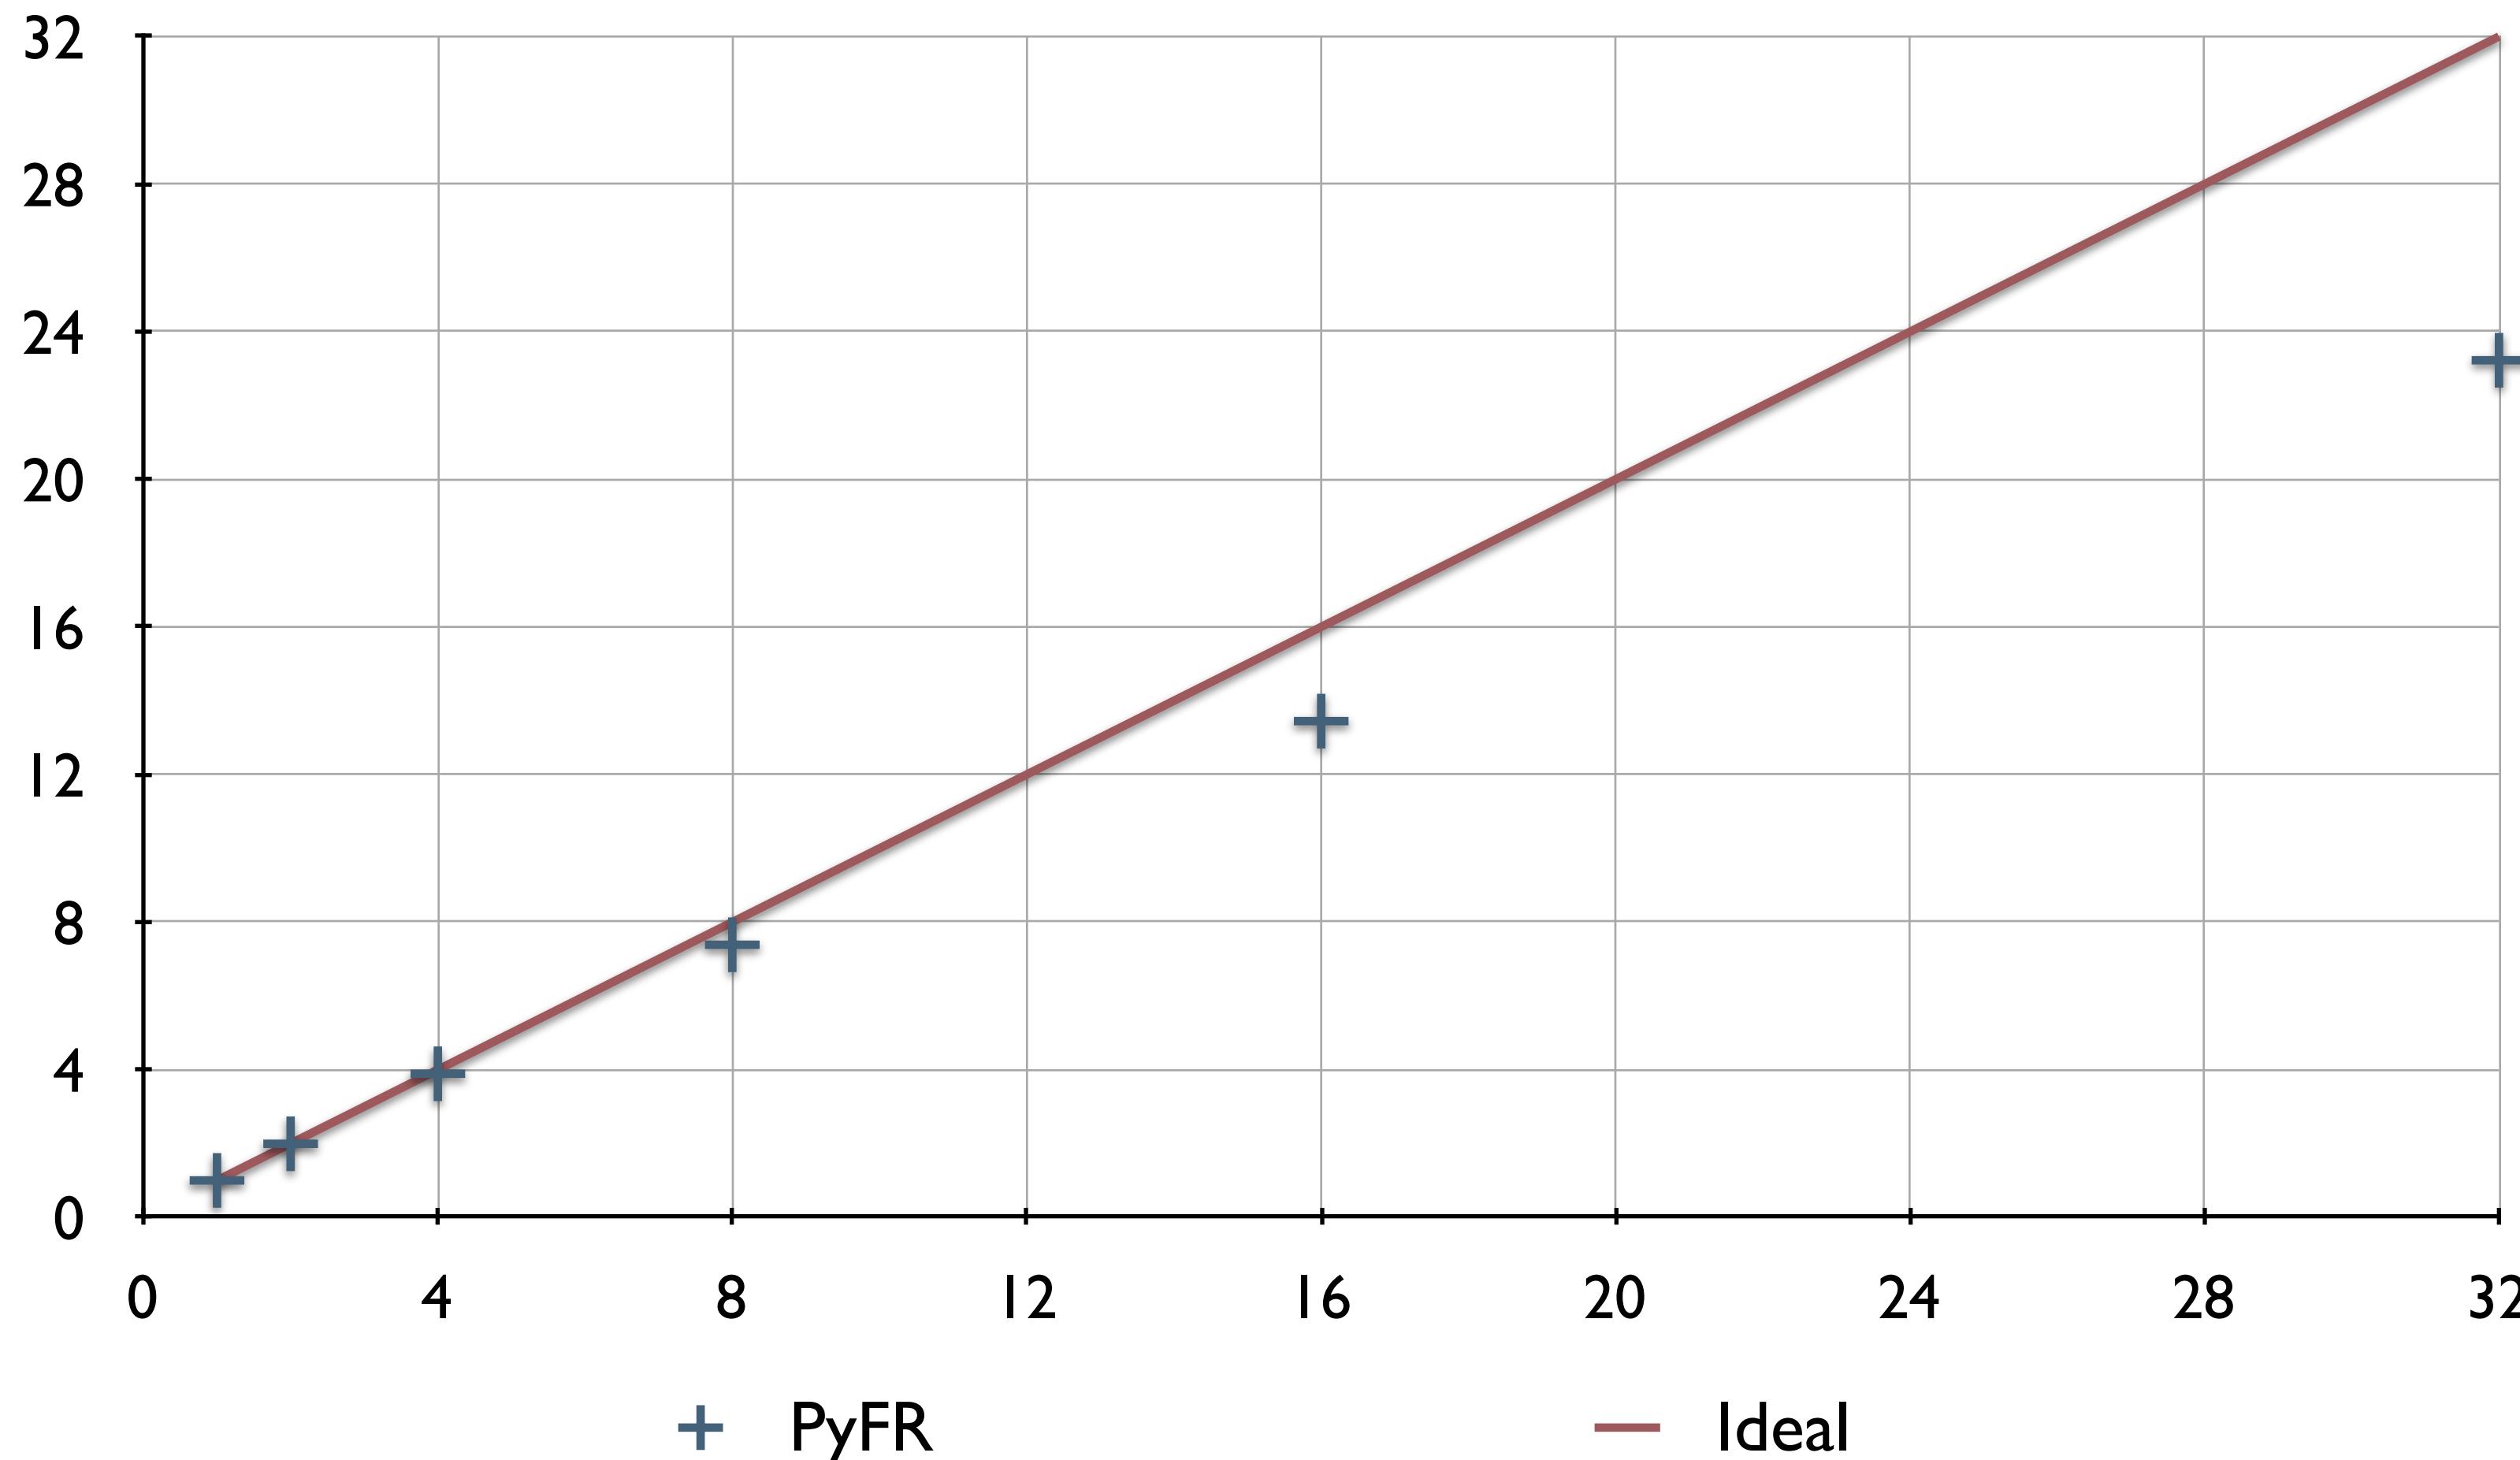
Results

- 3D Navier-Stokes weak scaling on Emerald



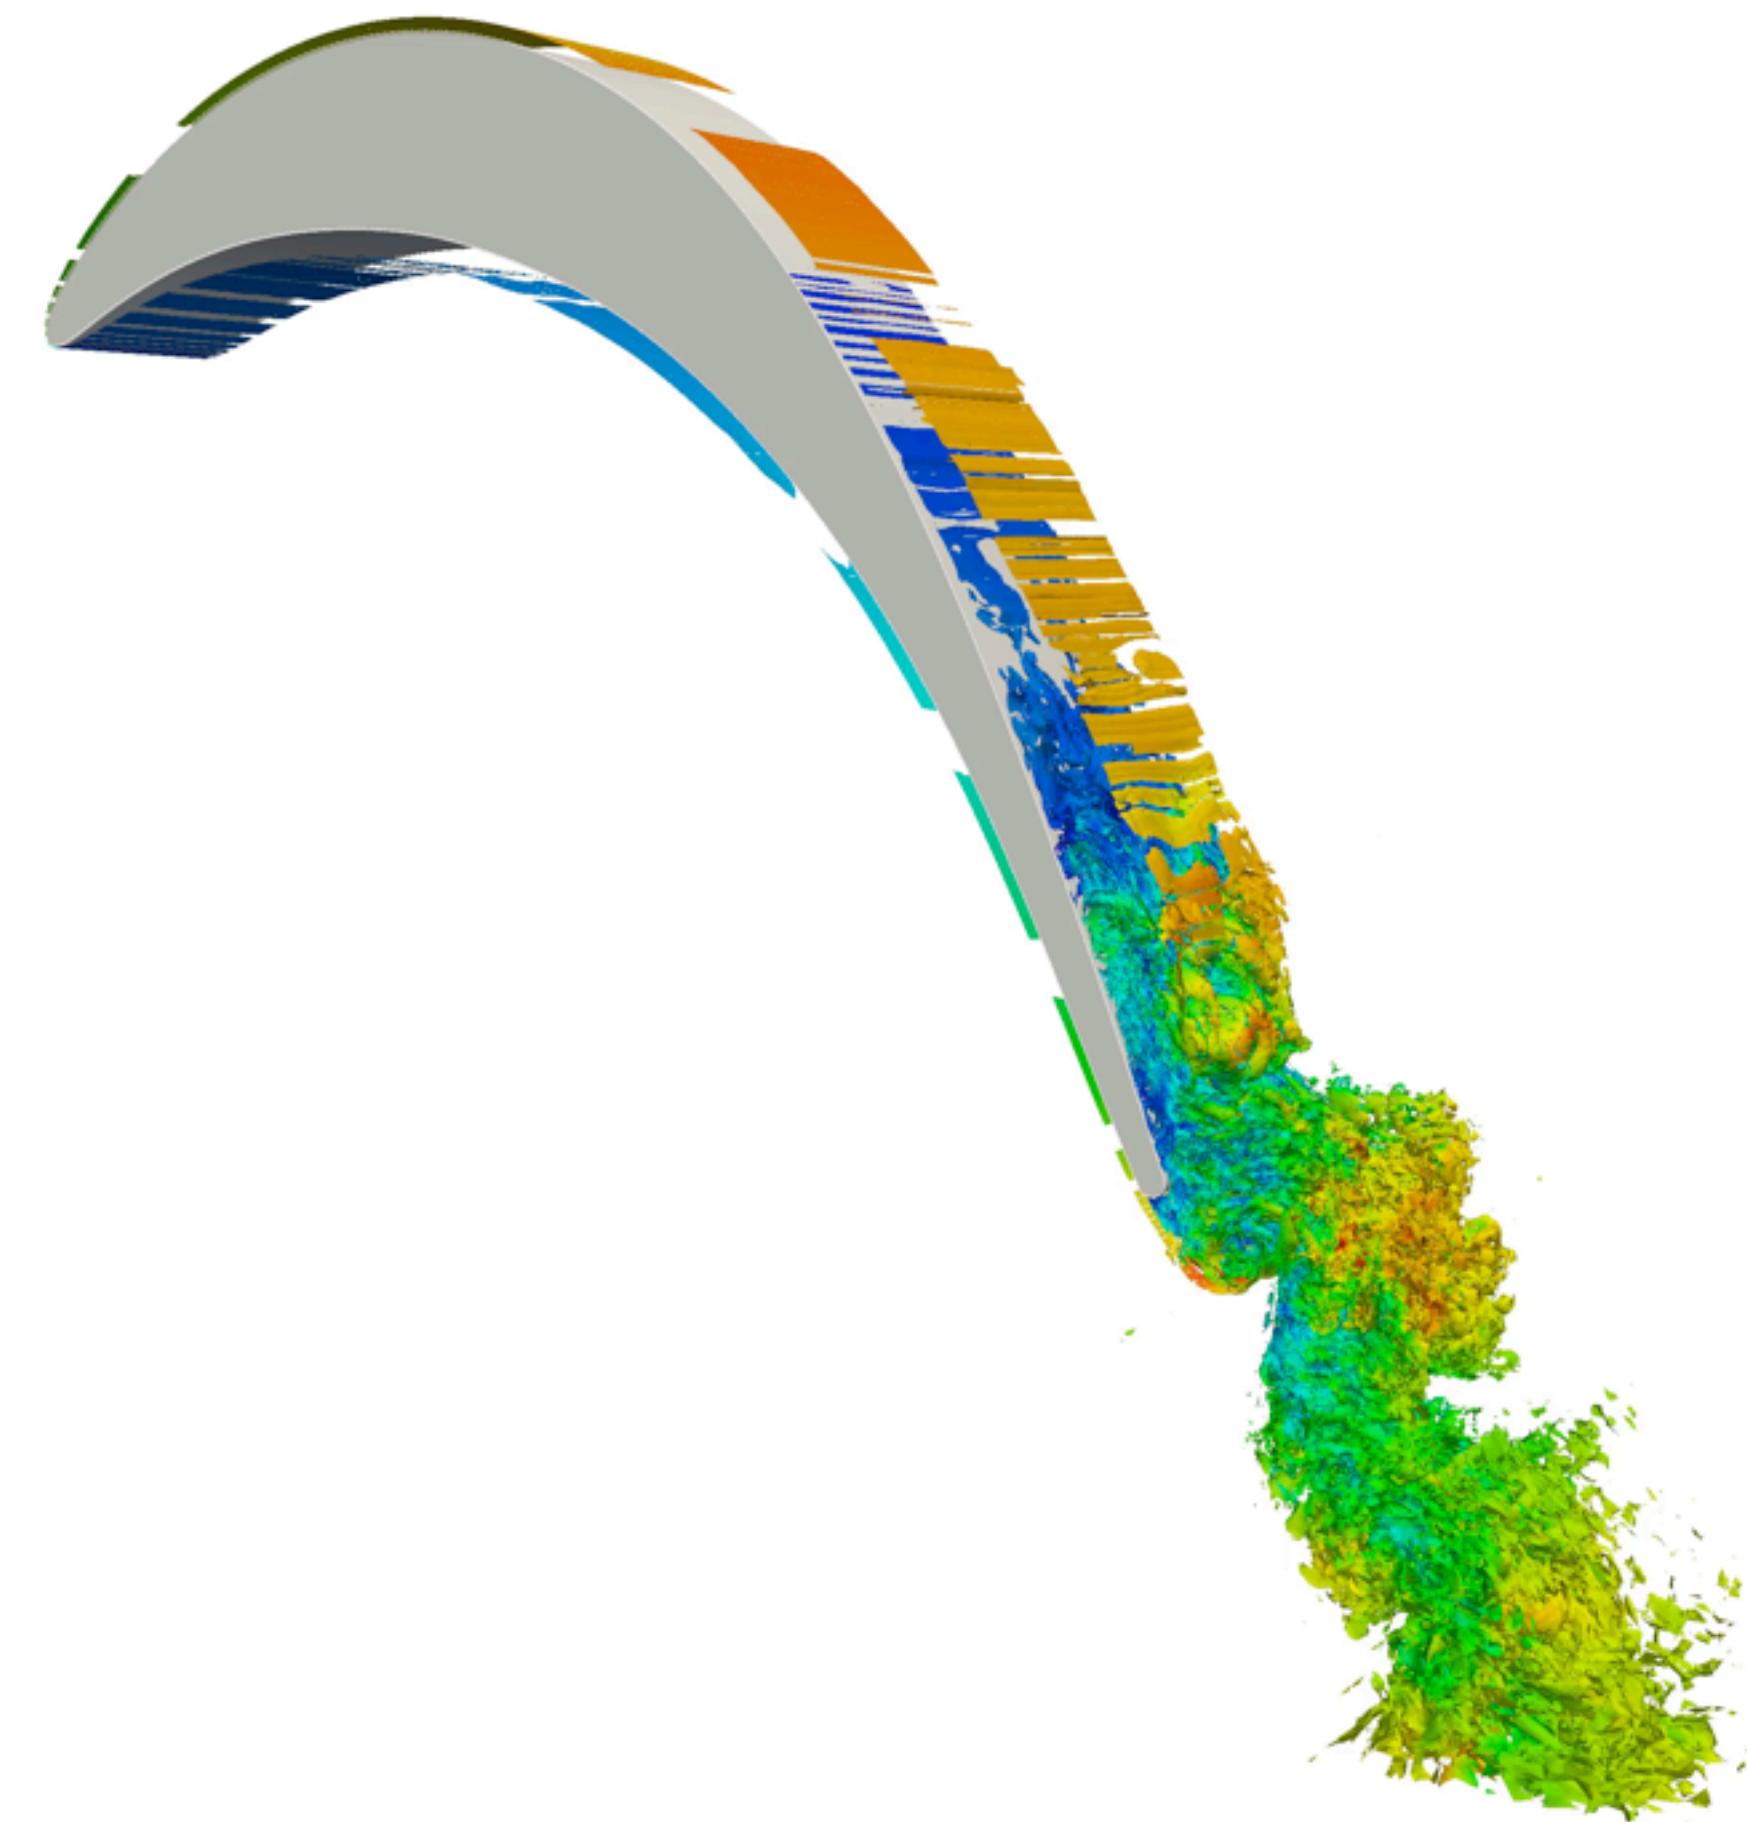
Results

- 3D Navier-Stokes **strong scaling** on Emerald



Results

- T106c Low Pressure Turbine Blade at $Re = 80,000$



Summary

- Interested in FR methods
- Interested in simple and efficient Python based implementations that target multiple hardware platforms
- Interested in application to solve real world industrial flow problems

Summary

- Funded and supported by



Engineering and Physical Sciences
Research Council



- Any questions?
- E-mail: freddie.witherden08@imperial.ac.uk